Breaking the Byte Barrier: Why 16-Bit Native Entropy Coding Matters for Medical Image Compression

Kuldeep Singh
Innovaccer Inc.
kuldeep.singh@innovaccer.com
https://github.com/pappuks/medical-image-codec

Abstract

The dominant entropy coders in modern compression systems—Huffman, arithmetic coding, LZ77, and FSE/ANS as used in Zstandard—were designed for 8-bit byte streams. Medical imaging breaks this assumption: DICOM stores pixels at 10–16 bits per sample, yielding symbol alphabets of 1,024 to 65,536 entries. We show that this mismatch is not merely inconvenient but structurally costly: byte-splitting a 16-bit residual distribution discards mutual information between the high and low bytes, inflating compressed output by 10–22% compared to a 16-bit-native entropy coder.

We present **MIC** (Medical Image Codec), a lossless codec that addresses this gap with four contributions:

1. **16-bit-native entropy coding**: An extended FSE implementation supporting up to 65,535 symbols, paired with a 16-bit run-length encoder that captures word-level structure invisible to byte-oriented compressors. MIC outperforms Delta+Zstandard by 10–22% on all tested modalities.

2. **Multi-state ANS decoding as an ILP design axis**: A four-state interleaved ANS decoder that breaks the serial dependency chain inherent in table-based ANS, with platform-specific BMI2 (amd64) and scalar (arm64) assembly kernels. We provide a formal latency model showing that the four-state design reduces amortised per-symbol latency from L to $L/4$ cycles, achieving 66–142% faster isolated FSE decompression with zero compression ratio penalty.

3. **The simplicity thesis**: Empirical evidence that for images where a simple spatial predictor achieves >90% of optimal decorrelation, additional transform complexity yields diminishing returns on lossless ratio while incurring proportional throughput cost. A ~1,100-line Delta+RLE+FSE pipeline matches or beats JPEG 2000's wavelet+EBCOT on 7/8 medical modalities—at dramatically higher throughput and a fraction of HTJ2K's implementation complexity.

4. **Browser-native decoding for ubiquitous distribution**: A pure JavaScript ES module (~20 KB, zero dependencies) and a Go WebAssembly build that decompress MIC files in any modern browser. PICS (Parallel Image Compressed Strips) extends this to multi-core browser decoding via `worker_threads`, achieving up to 483 MB/s on a 12-core workstation.

We evaluate MIC on 21 clinical DICOM datasets spanning MR, CT, CR, X-ray, mammography, nuclear medicine, radiography, secondary capture, and fluoroscopy (XA) using fair in-process benchmarking via CGO bindings against HTJ2K (OpenJPH), JPEG-LS (CharLS), and Delta+Zstandard. MIC achieves compression ratios of 1.7×–8.9×. MIC-4state-C exceeds HTJ2K decompression speed on 19/21 images single-threaded (ARM64) and 16/21 (AMD64, AVX2). PICS-C-8 exceeds HTJ2K on **all 21 images** on both platforms.

1. Introduction: The 16-Bit Alphabet Gap

Medical imaging generates enormous volumes of data. A single digital breast tomosynthesis (DBT) acquisition view produces 60 or more reconstructed slices—e.g., 69 frames at 2457 × 1890 pixels, 10-bit depth on the Hologic Selenia Dimensions—totalling over 600 MB of raw pixel data per view. A complete bilateral four-view screening study exceeds 2 GB uncompressed. Efficient lossless compression is essential for archival storage, network transmission, and real-time rendering in clinical PACS and diagnostic viewers.

The DICOM standard [1] supports several transfer syntaxes for lossless compression, including JPEG 2000 [2], JPEG-LS [4], and RLE Lossless. Each has well-documented tradeoffs between compression ratio and decompression speed [25, 26]. But all share a less-discussed limitation: **they were designed in an era of 8-bit data**.

JPEG's Huffman and arithmetic coders, JPEG-LS's Golomb-Rice coder, JPEG 2000's MQ coder, and Zstandard's FSE backend all treat each coding unit as a small fixed-size symbol (typically 4–8 bits). FSE as used in Zstandard [9] caps its alphabet at 4,096 symbols. Deflate, LZ4, Brotli, and even JPEG XL [5] operate on 8-bit byte streams.

Medical images break this assumption. DICOM stores pixels at 10, 12, or 16 bits per sample, yielding alphabet sizes of 1,024, 4,096, or 65,536 symbols. Even after spatial delta prediction, the residual alphabet for a 16-bit CT image can span tens of thousands of distinct values. An entropy coder that assumes an 8-bit alphabet must either:

- **Re-encode 16-bit symbols as byte pairs**, losing inter-byte correlation and doubling the number of coding steps, or

- **Truncate the alphabet** at 256 and fall back to raw storage for rare symbols.

This mismatch has measurable consequences. We show (Section 6) that applying Zstandard to delta-encoded 16-bit medical images loses 10–22% of achievable compression compared to a 16-bit-native approach. The loss is structural. A run of 1,000 identical 16-bit zero residuals is a single same-run record in 16-bit RLE, but 2,000 bytes with no obvious structure to a byte-level LZ77 matcher.

Contributions

1. We expose and quantify the 16-bit alphabet gap and demonstrate a 16-bit-native RLE+FSE pipeline that closes it (Section 3).
2. We formalise multi-state ANS decoding as an ILP design axis, providing a latency model and platform-specific assembly implementations that achieve 66–142% speedup with no ratio penalty (Section 4).
3. We present the simplicity thesis—empirical evidence that a simple spatial predictor with 16-bit entropy coding outperforms or matches wavelet+context-adaptive approaches at dramatically higher throughput (Section 6).
4. We demonstrate browser-native decoding as a distribution primitive: a ~20 KB JavaScript decoder and PICS achieve 483 MB/s in a browser (Section 5).

2. Background and Related Work

2.1 Lossless Image Compression in DICOM

JPEG 2000 [2] employs a 5/3 reversible integer wavelet transform with embedded block coding (EBCOT). It achieves excellent compression ratios but imposes significant decompression overhead. **High-Throughput JPEG 2000 (HTJ2K)** [3] replaces EBCOT with a faster block coder, substantially improving decode speed.

JPEG-LS [4] uses context-adaptive prediction followed by Golomb-Rice coding. It provides fast encoding with moderate compression ratios.

RLE Lossless (DICOM TS 1.2.840.10008.1.2.5) uses byte-level PackBits run-length encoding. It is simple but typically achieves <1.5× compression because it operates on bytes rather than 16-bit pixel values.

2.2 Entropy Coding and the Byte Assumption

Asymmetric numeral systems (ANS) [6, 7] generalise arithmetic coding into a single-state machine with integer state transitions. Finite State Entropy (FSE) [8] is a table-driven tANS implementation that achieves near-optimal compression with $O(1)$ per-symbol operations. FSE was popularised by Zstandard [9]. Recent theoretical analyses have formalised ANS efficiency bounds [12] and statistical interpretations [13].

All of these coders share one assumption: small alphabets. Huffman coding becomes impractical above ~4,096 symbols. FSE in Zstandard caps at 4,096. The entire ecosystem was built for bytes.

2.3 The Information-Theoretic Cost of Byte-Splitting

Consider a 16-bit symbol X decomposed into its high byte X_H and low byte X_L. The joint entropy satisfies:

$$H(X) = H(X_H, X_L) = H(X_H) + H(X_L \mid X_H). \tag{1}$$

A byte-level coder that processes X_H and X_L independently achieves at best:

$$H(X_H) + H(X_L) = H(X) + I(X_H; X_L), \tag{2}$$

where $I(X_H; X_L)$ is the mutual information between the two bytes. This mutual information is *discarded*—it represents correlation that the byte-level coder cannot exploit.

For medical image residuals after delta prediction, this mutual information is substantial. When a residual is small (say, ±30), the high byte is always `0x00` and carries essentially no information. On our test images, $I(X_H; X_L)$ of delta residuals ranges from 0.3 bits/symbol (MR) to 1.1 bits/symbol (MG3), corresponding to 8–22% of the total entropy—closely matching the 10–22% empirical advantage of MIC over Delta+Zstandard.

2.4 Color Transforms for RGB Modalities

The YCoCg-R transform [24] is a reversible integer approximation of RGB-to-YCbCr. It decorrelates RGB channels into luminance (Y) and chrominance (Co, Cg) with no loss of precision, and is used in MIC for ultrasound and visible-light images.

3. Pipeline Design: Delta + 16-Bit RLE + Extended FSE

MIC processes 16-bit greyscale pixels through three sequential stages. Each stage is simple enough to decode in a single sequential pass with minimal branching.

```
Raw 16-bit Pixels
   v pixel - avg(top, left)
Delta Encoding
   v same / diff runs
Run-Length Encoding (16-bit native)
   v table-driven ANS
FSE (tANS) Entropy Coding
   v
Compressed Bitstream
```

Figure 1: MIC compression pipeline.

3.1 Delta Encoding

The delta encoder computes the prediction residual for each pixel as:

$$r_{i,j} = x_{i,j} - \left\lfloor \frac{x_{i-1,j} + x_{i,j-1}}{2} \right\rfloor, \tag{3}$$

where $x_{i,j}$ is the pixel at row i, column j. Boundary pixels use only the available neighbour; the top-left corner pixel is stored raw.

For 16-bit images, residuals can span $[-65535, +65535]$. Rather than widening to 32 bits, MIC uses an **overflow delimiter** scheme derived from the image's effective bit depth $d = \lceil \log_2(\max(x) + 1) \rceil$:

- Residuals with $|r| < 2^{d-1} - 1$ are stored directly as uint16 values.
- Larger residuals are preceded by a delimiter value $2^d - 1$, followed by the raw pixel value.

Overflow delimiter: no collision by construction. In-range residuals are stored as `uint16(deltaThreshold + diff)`, mapping to the interval $[1, 2^d - 3]$. The delimiter is $D = 2^d - 1$. Since $2^d - 3 < 2^d - 1$, a direct residual can never equal D; the two ranges are disjoint.

Boundary-check-free interior loop. The delta decompressor uses four separate code paths for the corner pixel, first-row pixels, first-column pixels, and interior pixels. The interior loop handles $(w-1)(h-1)$ out of wh pixels without any per-pixel boundary checks, leaving only the highly-predictable overflow delimiter branch.

Why not the MED predictor? We implemented the JPEG-LS MED predictor through the full RLE+FSE pipeline. Result: geometric mean improvement of only +0.9% ratio, with MG4 actually regressing by −1.7%, at 1.5–2× slower decompression due to the three-way conditional branch and diagonal neighbour dependency.

3.2 Run-Length Encoding (16-Bit Native)

After delta coding, medical images produce long runs of identical residual values (especially zero). MIC's RLE stage encodes the symbol stream as alternating **same runs** and **diff runs**:

- **Same run**: a count c followed by a single value repeated c times.
- **Diff run**: a count c followed by c distinct values.

Why 16-bit RLE matters. A run of 1,000 identical 16-bit zero residuals is a single same-run record in MIC's RLE (2 uint16 words). The same data viewed as 2,000 bytes has no obvious run structure to a byte-level compressor.

The run headers use a midpoint protocol: counts ≤ `midCount` signal same runs; counts > `midCount` signal diff runs.

Same-run fast path. Same-value runs—the most common pattern after delta coding, often >80% of all runs on smooth medical images—are fast-pathed in the decoder to return the cached value without touching the compressed-data slice. Each output symbol in a same run requires only a counter decrement.

3.3 FSE Entropy Coding (Extended to 65,535 Symbols)

The final stage encodes the RLE output using tANS. **MIC extends FSE to support up to 65,535 distinct symbols** (versus 4,095 in Zstandard), which is necessary for 16-bit medical images where delta residuals can span a wide range.

Adaptive table sizing. MIC automatically adjusts the FSE table log via thresholds derived from *symbolLen* and *symbolDensity* = *inputLen*/*symbolLen*:

- Default: *tableLog* = 11
- Bumped to 12 when *symbolDensity* > 32 and *symbolLen* > 128
- Bumped to 13 when *symbolLen* > 512 and *symbolDensity* > 16

This provides 4–7% better compression on CR and mammography images (Table 1).

Table 1: Effect of adaptive tableLog (11→12).

Image	tableLog=11	Adaptive	Gain
CR	3.47×	3.63×	+4.4%
MG1	7.99×	8.57×	+7.1%
MG2	7.98×	8.55×	+7.1%

Dynamic working-set sizing. Encode and decode tables are sized to the actual symbol range rather than the theoretical 65,536 maximum. For 8-bit images this reduces the working set from 512 KB to approximately 2 KB, substantially improving cache utilisation.

3.4 Pipeline Selection Guidelines

Table 2: Recommended pipeline by use case.

Condition	Pipeline
Eff. bit depth ≤ 12, speed	Delta+RLE+FSE (4-state)
Eff. bit depth ≤ 12, ratio	Wavelet V2 SIMD
Full 16-bit range (CT)	Delta+RLE+FSE (avoid wavelet)
≥ 0.5 MB, multi-core avail.	Add PICS (2–8 strips)
Multi-frame sequence	MIC2 independent mode
Single-frame RGB (US, VL)	MICR (YCoCg-R, no tiling)
Lossy/progressive	Use HTJ2K (MIC is lossless only)

4. Multi-State ANS: ILP as a First-Class Design Axis

4.1 The Serial Dependency Problem

The standard FSE decode loop has a serial carried-dependency chain:

$$s[n+1] = \texttt{decTable}[s[n]].\mathit{newState} + \texttt{getBits}(\texttt{decTable}[s[n]].\mathit{nbBits})\,. \quad (4)$$

With a table-lookup latency of approximately $L \approx 4$ cycles (L1 cache hit), the loop is limited to roughly one symbol per 4 cycles, regardless of available execution units. Modern CPUs have 4–6 execution ports; the single-state decoder uses exactly **one** dependency chain, leaving 75–83% of the hardware capacity idle.

4.2 Multi-State Interleaved Decoding

MIC runs four independent state machines on interleaved symbol positions:

$$\begin{aligned}
\text{symbol}[0] &\leftarrow s_A, \quad \text{symbol}[4] \leftarrow s_A, \ \ldots \\
\text{symbol}[1] &\leftarrow s_B, \quad \text{symbol}[5] \leftarrow s_B, \ \ldots \\
\text{symbol}[2] &\leftarrow s_C, \quad \text{symbol}[6] \leftarrow s_C, \ \ldots \\
\text{symbol}[3] &\leftarrow s_D, \quad \text{symbol}[7] \leftarrow s_D, \ \ldots
\end{aligned}$$

Correctness. The encoder writes symbols at positions $0, 4, 8, \ldots$ into state A's bitstream; positions $1, 5, 9, \ldots$ into

state B's; etc. The decoder reverses this assignment. Each chain operates on a strict subset of positions with no data dependency on other chains; correctness follows from the correctness of single-state FSE applied independently to each subset.

Latency model.

Table 3: Amortised latency and speedup for k-state FSE decoder.

States	Amortised latency	Theoretical	Empirical
1	L cy/sym	1.0×	1.0×
2	$L/2$ cy/sym	2.0×	1.3×
4	$L/4$ cy/sym	4.0×	2.0×

Stream format. The compressed stream is prefixed with a magic sequence (`0xFF`, `0x02` for two-state; `0xFF`, `0x04` for four-state) followed by a 4-byte symbol count, enabling the auto-detect dispatcher to route streams transparently.

4.3 Platform-Specific Assembly Kernels

AMD64 (BMI2). Uses `SHLXQ`/`SHRXQ` for 3-operand variable shifts, enabling four independent bit-extraction sequences per iteration without CL register contention.

ARM64. Uses `LSLV`/`LSRV` variable-shift instructions which natively take the count from any general-purpose register.

AMD64 AVX2 — wavelet kernels. `wt53PredictAVX2` and `wt53UpdateAVX2` process 8 `int32` values per cycle using `VPADDD`/`VPSUBD`/`VPSRAD`, operating on contiguous arrays produced by the blocked column layout (8 columns per cache line).

Fallback. Pure-Go implementations of all routines provide identical functionality on non-amd64/arm64 targets (WebAssembly, RISC-V, etc.).

4.4 Isolated FSE Decompression Throughput

Table 4: Isolated FSE decompression throughput (MB/s) for 1-, 2-, and 4-state decoders (Intel Xeon @ 2.80 GHz). MB/s over uncompressed RLE symbol stream.

Image	1-state	2-state	4-state	4 vs. 1
MR	164	207	298	+82%
CT	113	126	195	+73%
CR	177	182	310	+75%
XR	193	172	320	+66%
MG1	158	207	380	+140%
MG3	576	890	1,343	+133%
MG4	256	321	620	+142%

5. Browser-Native Decoding

MIC ships two browser decoder implementations:

Table 5: Browser decoder options.

Decoder	Size	Deps.	Build
`mic-decoder.js`	~20 KB	None	None
`mic-decoder.wasm`	~2.5 MB	`wasm_exec.js`	`go build`

The JavaScript decoder is a complete, self-contained ES module with zero npm dependencies, working in any browser since 2020 (Chrome 67+, Firefox 68+, Safari 14+).

Single-threaded throughput (Apple M2 Max, Node.js v24.8):

Table 6: JavaScript decoder throughput (4-state), Node.js v24.8, M2 Max.

Image	Dimensions	Ratio	ms	MB/s
MR	256×256	2.35×	3.0	42
CT	512×512	2.24×	13.5	37
CR	1760×2140	3.69×	161.2	45
MG1	1996×2457	8.79×	80.9	116
MG3	3064×4774	2.29×	638.4	44

Parallel decoding via PICS + `worker_threads`. PICS files encode independent strips; each worker receives its strip, decompresses it independently, and returns the pixel buffer.

Table 7: PICS parallel JS decoder (`worker_threads`), Apple M2 Max.

Image	Strips	ms	MB/s	Speedup
MR (256^2)	4	1.3	94	2.57×
CT (512^2)	4	5.0	101	2.95×
CR (1760×2140)	8	31.7	227	5.53×
MG1 (1996×2457)	8	19.4	483	4.36×

MG1 (a 9.35 MB mammography image) decompresses in **19.4 ms** at 483 MB/s using 12 browser workers—well into real-time territory for diagnostic viewing.

Distribution model. A web-based DICOM viewer can fetch a `.mic` file directly from object storage (S3, GCS) and decode it client-side—zero server-side compute, zero transcoding latency, works offline, works in service workers. This is not possible today with HTJ2K or JPEG-LS.

6. Evaluation

6.1 Benchmarking Methodology and Dataset

Methodology. All comparisons use in-process CGO library calls for both MIC and competing codecs, measured on the same hardware in the same process (`BenchmarkAllCodecs`, `-benchtime=10x`).

Platforms. Apple M2 Max (12-core ARM64) and Intel Core Ultra 9 285K (24 P-core AMD64, `-march=native`).

Dataset. 21 clinical DICOM images spanning 10 modalities, drawn from publicly available de-identified datasets (NEMA WG-04 compsamples, NEMA 1997 CD, Clunie DBT Case 1). See Table 8.

6.2 Compression Ratios

Analysis. Mammography achieves the highest ratios (up to 8.87×) due to smooth tissue regions producing near-zero RLE runs. CT with full 16-bit dynamic range achieves 2.24× (MIC) vs. 1.67× (Wavelet): coefficient overflow in the 5/3 lifting step inflates compressed size. PICS improves ratio on large images

Table 8: Test dataset — 21 clinical DICOM images spanning 10 modalities.

Image	Modality	Dimensions	Raw (MB)
MR	Brain MRI	256×256	0.13
CT	CT scan	512×512	0.50
CR	Comp. radiography	2140×1760	7.18
XR	X-ray	2048×2577	10.1
MG1	Mammography	2457×1996	9.35
MG2	Mammography	2457×1996	9.35
MG3	Mammography	4774×3064	27.3
MG4	Mammography	4096×3328	26.0
CT1	CT scan	512×512	0.50
CT2	CT scan	512×512	0.50
MG-N	Mammography	3064×4664	27.3
MR1	Brain MRI	512×512	0.50
MR2	Brain MRI	1024×1024	2.00
MR3	Brain MRI	512×512	0.50
MR4	Brain MRI	512×512	0.50
NM1	Nuclear medicine	256×1024	0.50
RG1	Radiography	1841×1955	6.86
RG2	Radiography	1760×2140	7.18
RG3	Radiography	1760×1760	5.91
SC1	Sec. capture	2048×2487	9.71
XA1	Fluoroscopy (XA)	1024×1024	2.00

because each strip receives a dedicated FSE table that specialises to its local residual distribution: $\sum_k |S_k| \cdot H(S_k) < |S| \cdot H(S)$.

16-bit gap vs. Zstandard. MIC outperforms Delta+Zstandard (at level 19) by 10–22% on all tested modalities. The gap is structural: mutual information $I(X_H; X_L)$ of delta residuals ranges from 0.3 bits/symbol (MR) to 1.1 bits/symbol (MG3), matching the empirical gap exactly.

6.3 Decompression Speed — ARM64

MIC-4state-C is the recommended production variant. It is the fastest single-threaded decoder on **19 of 21 images on ARM64**. PICS-C-8 exceeds HTJ2K on **all 21 images**.

MIC-4state-C leads 13/21 images single-threaded. **MIC-4state-C is 1.7–5.0× faster than JPEG-LS on all 21 images.** PICS-C-8 peaks at 3,773 MB/s on MG2, enabling sub-millisecond decompression of large mammography images on a 12-core workstation.

6.4 Decompression Speed — AMD64

On AMD64 (`-O3 -march=native`), `MIC-4state-SIMD` activates the BMI2/PDEP decoder, and Wavelet+SIMD gains AVX2 kernels.

MIC-4state-SIMD leads 16/21 images single-threaded; HTJ2K leads on MG1, MG2, MG4, CT, and RG3. **PICS-C-8 exceeds HTJ2K on all 21 images** on both platforms.

6.5 Encoding Speed

On ARM64, MIC-4state-C achieves 243–644 MB/s single-threaded (1.2–2.5× faster than MIC-Go). PICS-8 scales to **1.2–2.1 GB/s** on large images using 8 goroutines, sufficient for real-time PACS ingestion. HTJ2K encode is comparable to MIC-Go on most images; JPEG-LS is consistently 2–4× slower to encode. Wavelet+SIMD encode is 2–4× slower than MIC-Go due to the multi-level forward transform cost.

6.6 Multi-Frame and Single-Frame RGB Compression

Multi-frame (MIC2) — 69-frame breast tomosynthesis (2457 × 1890 × 69, 614 MB raw):

Independent mode outperforms temporal mode on this dataset because tomosynthesis frames differ in X-ray projection angle, not temporal motion. Temporal mode is designed for genuine temporal sequences (cardiac cine MRI, fluoroscopy).

Single-Frame RGB (MICR) — YCoCg-R + Delta+RLE+FSE applied to the full image (NEMA compsamples):

7. Discussion

7.1 Signal Processing and Lossless Compression

For **lossy** compression, filter-bank decomposition is highly effective—high-frequency subbands can be quantised with little perceptual impact. For **lossless** compression the picture is more nuanced. The evaluation results (Section 6)—particularly the CT wavelet regression (1.67× vs. 2.24× for Delta)—are explained by three structural mechanisms:

1. **Coefficient overflow.** Even with integer lifting (the 5/3 Le Gall wavelet), the predict step can produce detail coefficients exceeding the input dynamic range. For 16-bit inputs, worst-case expansion per level is 3/2; for $L = 5$ levels, $(3/2)^5 \approx 7.6\times$—well beyond uint16, forcing int32 promotion or escape coding.

2. **The update step reintroduces correlation.** The lifting update adjusts low-pass coefficients to preserve the signal mean. For lossless coding of images already well-predicted by a simple linear filter, it adds residual energy to the low-pass band without a compensating reduction elsewhere.

3. **Higher memory bandwidth.** The 2D wavelet inverse transform requires two passes, each visiting every element at int32 width (4 bytes vs. 2 bytes per sample).

By contrast, simple spatial prediction leaves a single residual image whose statistics are homogeneous across the frame—ideal for a uniform 16-bit entropy coding scheme applied once.

7.2 The Cost-of-Complexity Frontier

JPEG-LS trades 58% of decompression speed for 10% more compression—unfavorable for clinical systems that decompress 10× more often than they compress. HTJ2K's ~20,000-line implementation is ~18× larger than MIC's Delta+RLE+FSE pipeline and has no practical browser decoder, making it unsuitable for client-side web deployment.

7.3 Limitations

- **No lossy compression**, progressive decoding, or region-of-interest coding for greyscale images. The wavelet pipeline provides a natural foundation for a future lossy mode.

- **Test dataset**: 21 images across 10 modalities from three public de-identified repositories. Large-scale benchmarks from public repositories (e.g., TCIA) would further strengthen generality claims.

Table 9: Lossless compression ratios — all codec variants, 21 images. Bold = best per row among non-JPEG-LS codecs.

Image	Raw (MB)	MIC	Wavelet	PICS-4	PICS-8	HTJ2K	JPEG-LS
MR	0.13	2.35×	2.38×	2.28×	2.21×	2.38×	**2.52**×
CT	0.50	**2.24**×	1.67×	2.15×	1.96×	1.77×	2.68×
CR	7.18	3.69×	3.81×	3.70×	3.71×	3.77×	**3.96**×
XR	10.1	1.74×	**1.76**×	1.75×	**1.76**×	1.67×	**1.76**×
MG1	9.35	8.79×	8.67×	8.84×	**8.87**×	8.25×	8.91×
MG2	9.35	8.77×	8.65×	8.83×	**8.85**×	8.24×	8.90×
MG3	27.3	2.24×	2.32×	2.31×	**2.34**×	2.22×	2.38×
MG4	26.0	3.47×	3.59×	3.59×	**3.62**×	3.51×	3.71×
CT1	0.50	**2.79**×	2.49×	2.54×	2.29×	2.70×	3.19×
CT2	0.50	**3.49**×	2.87×	3.11×	2.72×	3.29×	4.54×
MG-N	27.3	2.24×	2.32×	2.31×	**2.34**×	2.23×	2.38×
MR1	0.50	2.09×	2.14×	2.10×	2.08×	2.13×	**2.30**×
MR2	2.00	3.28×	3.34×	3.31×	3.31×	3.35×	**3.52**×
MR3	0.50	3.93×	4.09×	3.89×	3.84×	4.33×	**4.51**×
MR4	0.50	4.12×	4.18×	4.09×	4.03×	4.21×	**4.49**×
NM1	0.50	5.15×	5.02×	5.26×	**5.28**×	5.76×	6.28×
RG1	6.86	1.70×	1.70×	1.70×	1.69×	1.63×	**1.72**×
RG2	7.18	4.23×	4.32×	4.28×	4.30×	4.32×	**4.51**×
RG3	5.91	6.08×	6.82×	6.11×	6.12×	6.99×	**7.31**×
SC1	9.71	3.71×	3.70×	3.73×	3.74×	3.85×	**4.73**×
XA1	2.00	5.01×	4.94×	5.04×	**5.03**×	4.88×	5.39×

- **Temporal prediction** is underperforming on the one available clinical dataset and needs evaluation on cardiac cine MRI and fluoroscopy.
- **Benchmark confidence intervals**: run-to-run variance is <2% on M2 Max and <5% on AWS instances; formal confidence intervals are deferred to future work.

7.4 Pathway to Clinical Impact

MIC could be registered as a DICOM Private Transfer Syntax, allowing PACS vendors to adopt it without modifying the DICOM standard. The ~20 KB JavaScript decoder enables a direct storage-and-serve distribution model: compressed MIC files can be fetched from object storage and decoded client-side, eliminating server-side transcoding.

8. Conclusion

We have presented MIC, a lossless medical image codec built on a simple observation: the compression ecosystem has a 30-year blind spot for 16-bit data. By extending FSE to 65,535 symbols and pairing it with a 16-bit-native RLE stage, MIC outperforms byte-oriented Zstandard by 10–22% on all tested modalities. The four-state interleaved ANS decoder demonstrates that entropy coder design should target instruction-level parallelism width, not just coding efficiency, achieving 66–142% FSE speedup with no ratio penalty.

The simplicity thesis is validated empirically: a three-stage Delta+RLE+FSE pipeline (~1,100 lines of Go) matches or beats HTJ2K on 18/21 medical images for lossless compression ratio, at dramatically higher throughput.

Key results across 21 clinical DICOM images spanning 10 modalities:

- **Compression ratios**: 1.7×–8.9× (greyscale); 1.56×–6.24× (single-frame RGB US/VL).
- **vs. HTJ2K**: MIC-4state-C exceeds HTJ2K on 19/21 images single-threaded (ARM64); MIC-4state-SIMD exceeds on 16/21 (AMD64 AVX2); PICS-C-8 exceeds HTJ2K on **all 21 images** on both platforms.
- **vs. JPEG-LS**: 1.7–5.0× faster decompression at 1–30% lower ratio depending on modality.
- **vs. Delta+Zstandard**: 10–22% better compression on all modalities.
- **Browser decoder**: ~20 KB pure JS decoder; PICS + `worker_threads` achieves 483 MB/s (MG1) with no server-side transcoding.
- **Compact**: ~1,100 lines of Go — approximately 18× less than HTJ2K's ~20,000-line implementation.

MIC is open-source: https://github.com/pappuks/medical-image-codec.

References

[1] NEMA, "Digital Imaging and Communications in Medicine (DICOM)," https://www.dicomstandard.org/, 2024.

[2] D. S. Taubman and M. W. Marcellin, *JPEG2000: Image Compression Fundamentals, Standards and Practice*. Springer, 2002.

[3] D. S. Taubman, "High throughput JPEG 2000 (HTJ2K): Algorithm, performance evaluation, and potential," *Proc. SPIE*, vol. 11137, 2019.

[4] M. J. Weinberger, G. Seroussi, and G. Sapiro, "The LOCO-I lossless image compression algorithm: Principles and standardization into JPEG-LS," *IEEE Trans. Image Process.*, vol. 9, no. 8, pp. 1309–1324, 2000.

Table 10: Decompression throughput (MB/s), Apple M2 Max (ARM64), `-O3`. † = image too small for PICS-C-8; use PICS-C-4.

Image	MIC-Go	MIC-4s-C	Wvlt+SIMD	HTJ2K	JPEG-LS	PICS-C-2	PICS-C-4	PICS-C-8
MR	144	**348**	248	265	102	530	**710**†	482
CT	191	**356**	316	307	137	524	955	**1,092**
CR	296	524	**567**	367	153	867	1,635	**2,661**
XR	308	533	**627**	334	108	874	1,666	**3,025**
MG1	482	683	678	**810**	409	1,205	2,112	**3,656**
MG2	479	686	697	**790**	416	1,225	2,120	**3,773**
MG3	308	**531**	422	338	153	864	1,673	**3,117**
MG4	417	**625**	516	548	184	1,093	2,004	**3,689**
CT1	239	**436**	425	362	182	686	1,013	**1,183**
CT2	238	439	**481**	375	175	676	1,041	**1,189**
MG-N	316	**536**	468	340	153	883	1,711	**3,175**
MR1	278	**521**	435	325	116	751	1,207	**1,402**
MR2	333	**563**	498	388	172	913	1,552	**2,466**
MR3	375	**639**	507	441	236	908	1,430	**1,614**
MR4	316	**571**	479	406	197	818	1,341	**1,558**
NM1	327	**632**	575	410	210	888	1,400	**1,679**
RG1	235	406	**584**	332	104	602	1,128	**2,017**
RG2	367	590	**644**	443	193	986	1,803	**3,194**
RG3	374	604	**656**	562	246	1,035	1,944	**3,302**
SC1	375	**587**	388	401	229	1,017	1,861	**3,279**
XA1	331	**576**	459	419	204	928	1,583	**2,493**

[5] J. Alakuijala *et al.*, "JPEG XL next-generation image compression architecture and coding tools," *Proc. SPIE 11137*, Sep. 2019.

[6] J. Duda, "Asymmetric numeral systems," arXiv:0902.0271, 2009.

[7] J. Duda, "Asymmetric numeral systems: entropy coding combining speed of Huffman coding with compression rate of arithmetic coding," arXiv:1311.2540, 2013.

[8] Y. Collet, "Finite State Entropy — A new breed of entropy coder," https://github.com/Cyan4973/FiniteStateEntropy, 2013.

[9] Y. Collet and M. Kucherawy, "Zstandard Compression and the 'application/zstd' Media Type," RFC 8878, IETF, Feb. 2021.

[10] E. S. Schwartz and B. Kallick, "Generating a canonical prefix encoding," *Commun. ACM*, vol. 7, no. 3, pp. 166–169, 1964.

[11] S. Mahapatra and K. Singh, "An FPGA-based implementation of multi-alphabet arithmetic coding," *IEEE Trans. Circuits Syst. I*, vol. 54, no. 8, pp. 1678–1686, 2007.

[12] D. Kosolobov, "Efficiency of ANS Entropy Encoders," arXiv:2201.02514, 2022.

[13] R. Bamler, "Understanding Entropy Coding With Asymmetric Numeral Systems (ANS): a Statistician's Perspective," arXiv:2201.01741, 2022.

[14] F. Giesen, "Interleaved entropy coders," arXiv:1402.3392, 2014.

[15] F. Giesen, "ryg_rans: Public domain rANS encoder/decoder," https://github.com/rygorous/ryg_rans, 2014.

[16] F. Lin, K. Arunruangsirilert, H. Sun, and J. Katto, "Recoil: Parallel rANS Decoding with Decoder-Adaptive Scalability," *Proc. 52nd ICPP '23*, pp. 31–40, 2023.

[17] A. Weissenberger and B. Schmidt, "Massively Parallel ANS Decoding on GPUs," *Proc. 48th ICPP '19*, Article 100, 2019.

[18] X. Wu and N. Memon, "Context-based, adaptive, lossless image coding," *IEEE Trans. Commun.*, vol. 45, no. 4, pp. 437–444, 1997.

[19] J. Sneyers and P. Wuille, "FLIF: Free Lossless Image Format based on MANIAC Compression," *Proc. IEEE ICIP*, 2016.

[20] D. Le Gall and A. Tabatabai, "Sub-band coding of digital images using symmetric short kernel filters and arithmetic coding techniques," *Proc. IEEE ICASSP*, pp. 761–764, 1988.

[21] W. Sweldens, "The Lifting Scheme: A custom-design construction of biorthogonal wavelets," *Appl. Comput. Harmon. Anal.*, vol. 3, no. 2, pp. 186–200, 1996.

[22] I. Daubechies and W. Sweldens, "Factoring wavelet transforms into lifting steps," *J. Fourier Anal. Appl.*, vol. 4, no. 3, pp. 247–269, 1998.

[23] A. R. Calderbank, I. Daubechies, W. Sweldens, and B.-L. Yeo, "Wavelet transforms that map integers to integers," *Appl. Comput. Harmon. Anal.*, vol. 5, no. 3, pp. 332–369, 1998.

Table 11: Decompression throughput (MB/s), Intel Core Ultra 9 285K (AMD64), `-O3 -march=native`. † = image too small for PICS-C-8.

Image	MIC-Go	MIC-4s-C	MIC-4s-SIMD	Wvlt+SIMD	HTJ2K	PICS-C-4	PICS-C-8
MR	251	501	**708**	194	570	301	124†
CT	231	403	487	364	**544**	**676**	339
CR	324	599	**744**	723	708	1,738	**2,435**
XR	363	601	**803**	681	570	1,714	**1,994**
MG1	617	789	1,119	746	**1,235**	1,872	**2,514**
MG2	562	800	1,166	848	**1,297**	**2,244**	2,085
MG3	357	633	669	**678**	644	1,823	**2,538**
MG4	523	743	773	**808**	916	2,124	**2,707**
CT1	322	520	**676**	525	657	**857**	423
CT2	293	514	636	**705**	627	**847**	687
MG-N	368	635	669	**705**	643	1,872	**2,191**
MR1	356	609	**766**	532	654	**945**	366
MR2	352	658	**975**	601	749	**1,121**	793
MR3	450	728	**809**	676	802	**920**	705
MR4	390	596	**861**	738	660	493	**701**
NM1	367	**717**	668	715	627	**783**	405
RG1	302	497	593	**705**	557	1,248	**1,494**
RG2	433	676	**861**	811	823	1,841	**1,912**
RG3	460	706	858	881	**894**	**1,969**	1,613
SC1	464	695	**888**	710	728	1,819	**1,889**
XA1	413	693	**836**	538	797	1,173	**1,513**

Table 12: 69-frame breast tomosynthesis compression.

Mode	Compressed	Ratio
Independent (spatial)	46.1 MB	**13.3**×
Temporal (inter-frame)	47.5 MB	12.9×

Table 13: Single-frame RGB compression ratios (MICR, NEMA compsamples).

Image	Dimensions	Ratio	Notes
US1	640×480	6.24×	Uniform bg
VL1–3	756×486	3.2–3.5×	Visible light
VL4	2226×1868	1.86×	High detail
VL5	2670×3340	1.56×	Skin texture
VL6	756×486	1.93×	—

Table 14: Pareto frontier: ratio vs. decompression throughput (approx. geo. means).

Codec	Ratio	Decomp.	LOC	Browser
Δ+RLE+FSE (Go)	3.12×	310 MB/s	1.1k	**Yes (20 KB)**
Δ+RLE+FSE (4s-C)	3.12×	530 MB/s	1.5k	**Yes**
Wavelet V2 SIMD	3.28×	780 MB/s	2.0k	Possible
HTJ2K (OpenJPH)	3.15×	460 MB/s	20k	No
JPEG-LS (CharLS)	3.44×	130 MB/s	5.0k	No
Δ+zstd-19	2.72×	300 MB/s	N/A	Partial

[24] H. S. Malvar and G. J. Sullivan, "YCoCg-R: A color space with RGB reversibility and low dynamic range," JVT Doc. JVT-I014, 2003.

[25] F. Liu *et al.*, "The Current Role of Image Compression Standards in Medical Imaging," *Information*, vol. 8, no. 4, p. 131, Oct. 2017.

[26] D. A. Clunie, "Lossless Compression of Grayscale Medical Images: Effectiveness of Traditional and State-of-the-Art Approaches," *Proc. SPIE Medical Imaging*, 2000.

[27] D. Taubman, A. Naman, M. Smith, P.-A. Lemieux, H. Saadat, O. Watanabe, and R. Mathew, "High Throughput JPEG 2000 for Video Content Production and Delivery Over IP Networks," *Front. Signal Process.*, vol. 2, Article 885644, Apr. 2022.

[28] S. Williams, A. Waterman, and D. Patterson, "Roofline: An insightful visual performance model for multicore architectures," *Commun. ACM*, vol. 52, no. 4, pp. 65–76, Apr. 2009.

[29] A. N. Aous, "OpenJPH: An open-source implementation of High-Throughput JPEG 2000," https://github.com/aous72/OpenJPH, 2024.

[30] Team CharLS, "CharLS: A C/C++ JPEG-LS library implementation," https://github.com/team-charls/charls, 2024.

[31] As of the submission date of this paper, no production-ready WebAssembly or JavaScript decoder exists for High-Throughput JPEG 2000. The OpenJPH repository [29] does not provide a WASM build or npm package.

[32] N. Tindall *et al.*, "fzstd: Pure JavaScript Zstandard decompressor," https://github.com/101arrowz/fzstd, 2024.