

A 16-Bit-Native Entropy Coding Pipeline for Lossless Medical Image Compression

Kuldeep Singh

Abstract—Lossless compression is essential in medical imaging, where archival storage and diagnostic workflows require exact pixel reconstruction from 10–16-bit-per-sample data. Most entropy coders target 8-bit byte streams; medical images break this assumption by producing residual alphabets of thousands of symbols after spatial prediction. This paper presents MIC, a lossless codec for Digital Imaging and Communications in Medicine (DICOM) images that operates natively on 16-bit residuals. MIC combines spatial prediction, 16-bit run-length encoding, and large-alphabet Finite State Entropy (FSE), a table-driven implementation of asymmetric numeral systems (ANS). Rather than splitting 16-bit pixels into byte pairs before entropy coding, MIC treats each predicted residual as a single 16-bit symbol, capturing correlations between the high and low byte directly—a structural advantage quantified at 0.3–1.1 bits per symbol on the test dataset. The implementation extends FSE to support up to 65,535-symbol alphabets and includes a multi-state interleaved decoder that exploits instruction-level parallelism (ILP) during decompression. Evaluation on 21 de-identified DICOM images spanning nine modalities shows the 16-bit-native pipeline improves compression ratio by 5–22% (geometric mean 14%) over a Delta+Zstandard baseline, reaching a geometric-mean ratio of 3.12 times—within 1% of High-Throughput JPEG 2000 (HTJ2K, 3.15 times). A multi-state interleaved decoder yields a 46% geometric-mean FSE-stage decompression speedup on ARM64 at four states and a further 7% at eight states. Single-threaded, MIC’s C decoder is faster than HTJ2K on every ARM64 image and on every AMD64 image we tested; a strip-parallel mode adds near-linear scaling for large images. A 20-kilobyte pure JavaScript decoder and a Go WebAssembly build enable client-side in-browser decoding, confirming that 16-bit-native entropy coding is a practical design point where throughput, simplicity, and deployment portability matter.

Index Terms—medical image compression, lossless compression, entropy coding, asymmetric numeral systems (ANS), finite state entropy (FSE), DICOM, high-bit-depth imaging.

I. INTRODUCTION

MEDICAL imaging systems routinely generate large volumes of high-bit-depth data. Modalities such as computed tomography (CT), magnetic resonance imaging (MR), digital radiography (computed radiography, CR; plain X-ray, XR), mammography, and tomosynthesis commonly store pixel values at 10, 12, or 16 bits per sample. A single digital breast tomosynthesis (DBT) acquisition view produces 60 or more reconstructed slices, totalling over 600 MB of raw pixel data per view. A complete bilateral four-view screening study exceeds 2 GB uncompressed. Efficient lossless compression is essential for archival storage, network transmission, and real-

time rendering in clinical Picture Archiving and Communication Systems (PACS) and diagnostic viewers.

DICOM [1] supports several established lossless transfer syntaxes, including JPEG 2000, HTJ2K, JPEG-LS (a lossless/near-lossless predictive standard), and Run-Length Encoding (RLE) Lossless. These methods offer different tradeoffs among compression ratio, decoding speed, and implementation complexity [24], [25]. JPEG 2000 and HTJ2K provide strong compression performance but rely on transform and block-coding pipelines that are comparatively complex. JPEG-LS uses predictive coding and often yields strong lossless ratios, but its throughput characteristics remain application-dependent. RLE Lossless is simple but typically provides limited compression because it operates on byte-oriented runs and does not include an effective decorrelation stage.

This work studies a different design point for lossless medical image compression: direct coding of **predicted 16-bit residual streams**. Throughout this paper, a *residual stream* denotes the sequence of per-pixel prediction errors $r_{i,j} = x_{i,j} - \hat{x}_{i,j}$ produced by a spatial predictor, kept at the native 16-bit precision of the source pixels rather than re-serialised into bytes. As a concrete example: if the predicted pixel value is 1020 and the true pixel is 1023, the residual is +3. Smooth regions therefore produce many residuals near zero, and these near-zero residuals recur frequently—properties that motivate a 16-bit-native run-length and entropy coding stage.

The central observation is that the dominant entropy coders in modern compression systems, including Huffman coding, arithmetic coding, Lempel-Ziv 1977 (LZ77), and FSE/ANS as used in Zstandard, were designed primarily for 8-bit byte streams. Medical images break this assumption: DICOM stores pixels at 10–16 bits per sample, yielding symbol alphabets of 1,024 to 65,536 entries. Even after spatial delta prediction, the residual alphabet for a 16-bit CT image can span tens of thousands of distinct values. An entropy coder that assumes an 8-bit alphabet must either re-encode 16-bit symbols as byte pairs (losing inter-byte correlation and doubling the number of coding steps) or truncate the alphabet at 256 and fall back to raw storage for rare symbols. This mismatch is structurally costly: for example, small signed residuals such as $-2, -1, 0, +1$ have highly predictable high bytes, so encoding the high and low bytes independently forces the coder to spend bits on information already implied by the other byte. On our test images, the mutual information $I(X_H; X_L)$ between the high and low bytes of delta residuals ranges from 0.3 bits/symbol (MR) to 1.1 bits/symbol (MG3), corresponding to 8–22% of total entropy and explaining MIC’s 5–22% empirical advantage over Delta + Zstandard.

Based on this observation, we present MIC, a lossless codec that combines:

- 1) a simple spatial predictor,
- 2) a 16-bit-native run-length representation of the residual stream, and
- 3) an extended FSE/tANS (table-based ANS) entropy coder supporting large active symbol alphabets.

The codec also includes multi-state interleaved entropy decoding to improve decompression throughput by increasing ILP.

A. Contributions

The main contributions of this paper are as follows.

- 1) **A 16-bit-native lossless coding pipeline for medical images.** We present a Delta + 16-bit RLE + extended FSE pipeline for grayscale medical image compression that consistently outperforms Delta + Zstandard on all 21 tested images, with the improvement attributable to the structural cost of byte-splitting 16-bit residuals.
- 2) **An extended table-based entropy coder for large residual alphabets.** The proposed implementation supports up to 65,535 distinct symbols. One 16-bit code value ($2^{16} - 1 = 65535$) is reserved as the overflow delimiter for full-range images, leaving up to 65,535 codes available for entropy coding. Tables are dynamically sized to the observed symbol range, reducing working-set size from 512 KB to approximately 2 KB for 8-bit images.
- 3) **A multi-state interleaved entropy decoder.** We implement and evaluate two-, four- and eight-state FSE decoding to exploit ILP. Four states deliver a geometric-mean +46% FSE-stage speedup on ARM64 (+8% on AMD64) and eight states deliver a further +7% on ARM64 (+2% on AMD64), with no compression ratio penalty.
- 4) **An empirical evaluation against standard codecs and a general-purpose baseline.** We compare MIC with HTJ2K (OpenJPH) [29], JPEG-LS (CharLS) [30], and Delta + Zstandard on a dataset of 21 DICOM images spanning nine modalities (described in Section V).
- 5) **A compact JavaScript/WebAssembly (WASM) decoder implementation.** We provide a ~20 KB pure JavaScript (JS) decoder and a Go WebAssembly build to explore deployment feasibility for client-side viewing.

B. Scope

This paper is scoped to **lossless single-frame grayscale medical image compression**. Multi-frame and RGB extensions are outside the scope of this work. The present work does not address lossy coding, progressive decoding, or region-of-interest coding.

C. Organization

Section II reviews related work and positions MIC relative to existing methods. Section III describes the proposed coding pipeline. Section IV presents the multi-state entropy decoder. Section V describes experimental methodology. Section VI reports compression and speed results. Section VII concludes.

II. RELATED WORK

A. Lossless Compression in Medical Imaging

JPEG 2000 [2] employs the reversible 5/3 wavelet transform [20], [23] with Embedded Block Coding with Optimized Truncation (EBCOT), achieving excellent ratios at significant decompression overhead. HTJ2K [3], [27] replaces EBCOT with a faster block coder, substantially improving decode speed. JPEG-LS [4], [5] uses context-adaptive Median Edge Detector (MED) prediction with Golomb-Rice coding; it is an important lossless baseline as a standardized predictive codec widely deployed in medical PACS. RLE Lossless (DICOM TS 1.2.840.10008.1.2.5) applies byte-level PackBits without decorrelation, yielding poor ratios (typically $<1.5\times$). Higher-complexity approaches such as the Context-based Adaptive Lossless Image Codec (CALIC) [18] and the Free Lossless Image Format (FLIF) [19] achieve strong ratios on natural images but lack DICOM standardization and browser decoders.

These codecs are the primary baselines for this work. The goal is not to supersede them but to investigate a simpler residual-domain design specifically targeting high-bit-depth images with strong throughput and deployment portability.

B. Entropy Coding and the Byte Assumption

Asymmetric numeral systems (ANS) [7], [8] replace arithmetic coding's interval subdivision with integer state transitions. Finite State Entropy (FSE) [9] is a table-driven tANS implementation with $O(1)$ per-symbol operations, popularized by Zstandard [10]. Recent work has formalized ANS efficiency bounds [12] and statistical interpretations [13]. In practice, Huffman coding becomes impractical above ~4,096 symbols and FSE in Zstandard is capped at 4,096; modern formats such as JPEG XL [6] likewise target byte streams.

Information-theoretic cost of byte-splitting. For a 16-bit symbol X decomposed into high byte X_H and low byte X_L , a byte-level coder achieves at best $H(X_H) + H(X_L) = H(X) + I(X_H; X_L)$. After delta prediction, when a residual is small (for example, ± 30), X_H is nearly deterministic given X_L , but it must still be encoded independently. On our test images, $I(X_H; X_L)$ ranges from 0.3 bits/symbol (MR) to 1.1 bits/symbol (MG3), corresponding to 8–22% of total entropy and closely matching MIC's empirical advantage over Delta + Zstandard.

C. ANS and Interleaved Decoding

Prior work has shown that interleaving or multi-stream organization can improve ANS decoder throughput. Giesen [14], [15] demonstrated the value of multi-stream range ANS (rANS); Lin *et al.* [16] extended parallel rANS to decoder-adaptive scalability; Weissenberger and Schmidt [17] showed massively parallel ANS decoding on GPUs. The present work does not claim novelty in ANS itself or in interleaving as a general concept. Rather, the contribution lies in applying these ideas in a 16-bit-native medical image codec, together with a specific large-alphabet implementation and throughput-oriented software design.

D. Positioning of This Work

Table I distinguishes MIC from prior work across five dimensions relevant to medical image compression.

TABLE I: Feature comparison of MIC with related codecs and entropy coding methods.

Feature	JPEG-LS	HTJ2K	Zstd	rANS	MIC
Lossless medical	Yes	Yes	General	General	Yes
Native 16-bit	No	No	No	Configurable	Yes
Multi-state ILP	No	No	No	Yes	Yes
Browser decoder	No	No	Partial	No	Yes
Source lines	5k	20k	N/A	N/A	1.1k

III. PROPOSED METHOD

Each stage of the pipeline described in Section I is simple enough to decode in a single sequential pass with minimal branching, as illustrated in Fig. 1.

A. Spatial Prediction

Let $x_{i,j}$ denote the pixel at row i , column j . For interior pixels, MIC predicts the current pixel from the average of the top and left neighbors:

$$\hat{x}_{i,j} = \left\lfloor \frac{x_{i-1,j} + x_{i,j-1}}{2} \right\rfloor \quad (1)$$

where the division is integer division rounded down (matching the decoder exactly), and all arithmetic is performed in unsigned integers. The residual is formed as:

$$r_{i,j} = x_{i,j} - \hat{x}_{i,j}. \quad (2)$$

Boundary pixels use only the available neighbor, and the first pixel is stored directly. This predictor was selected because it is computationally simple, branch-light in decoding, and effective on the current dataset.

Predictor comparison. We implemented and compared four predictors through the full RLE + FSE pipeline on the 21-image dataset: left-only (left neighbor, no top), average (MIC default), Paeth (PNG predictor [26]), and MED (JPEG-LS [4]). Table II summarizes geometric means and relative gains.

TABLE II: Predictor ablation summary: geometric means and win counts across all 21 images.

Predictor	Geo. Mean	Wins	Avg→X
Left-only	3.38×	3/21	−2.3%
Average (MIC)	3.46×	13/21	baseline
Paeth	3.48×	2/21	+0.5%
MED (JPEG-LS)	3.52×	16/21	+1.6%

MED decompression is 1.5–2× slower due to the three-way conditional branch and diagonal neighbor dependency that prevents the branch-free interior loop used by the average predictor. MED yields the best geomean ratio (+1.6% over Avg) but at significant speed cost; the average predictor is retained as the default for its consistent per-modality performance and branch-free decompression path.

B. Effective Bit Depth and Overflow Coding

For 16-bit images, residuals can span $[-65535, +65535]$. Rather than widening the main residual stream to 32 bits, MIC uses an overflow delimiter scheme derived from the effective image bit depth:

$$d = \lceil \log_2(\max(x) + 1) \rceil. \quad (3)$$

Residuals within an in-range interval are mapped directly to 16-bit codes. Residuals outside that interval are encoded using a delimiter followed by the raw pixel value.

Encoding procedure:

- 1) Compute $\text{deltaThreshold} = 2^{d-1} - 1$.
- 2) Compute delimiter $D = 2^d - 1$.
- 3) For each residual $r_{i,j}$: if $|r_{i,j}| < \text{deltaThreshold}$, emit $\text{uint16}(\text{deltaThreshold} + r_{i,j})$; otherwise emit delimiter D , then emit raw pixel $x_{i,j}$.

The threshold condition is strict ($<$, not $\leq$), so residuals with $|r| = \text{deltaThreshold}$ take the overflow path. Residual zero maps to code value deltaThreshold exactly. Table III illustrates the mapping for a representative bit depth.

TABLE III: Overflow coding example for $d = 12$ bits ($\text{deltaThreshold} = 2047$, $D = 4095$).

Residual r	Emitted code(s)	Notes
−2	2045	in-range
−1	2046	in-range
0	2047	in-range (= threshold)
+1	2048	in-range
+2	2049	in-range
±2047	4095, raw pixel	overflow

Non-collision proof. In-range residuals satisfy $|r| \leq \text{deltaThreshold} - 1$, so stored values range from 1 to $2^d - 3$. The delimiter is $D = 2^d - 1$. Since $2^d - 3 < 2^d - 1$, the stored value of a direct residual can never equal D ; the two ranges are disjoint.

Decoding procedure: Read the next uint16 value c . If $c \neq D$, recover the residual as $r = \text{int16}(c - \text{deltaThreshold})$ and reconstruct $x_{i,j} = \hat{x}_{i,j} + r$. If $c = D$, the raw pixel follows: read the next uint16 as $x_{i,j}$ directly.

C. 16-Bit Run-Length Encoding

The predicted residual stream is converted into an alternating sequence of:

- **Same runs:** a count and a repeated 16-bit value;
- **Diff runs:** a count followed by a sequence of distinct 16-bit values.

A same run is emitted only when at least three consecutive identical values are observed.

Midpoint protocol. Run headers are stored as uint16 values. A midpoint constant $\text{midCount} = 32767$ partitions the header space: a count $c \leq \text{midCount}$ signals a same run of length c ; a count $c > \text{midCount}$ signals a diff run of length $c - \text{midCount}$. The maximum same-run or diff-run length is therefore 32767 symbols; longer sequences are split into multiple headers transparently. The decoder distinguishes the two run types solely by comparing the header value against midCount .

MIC Compression Pipeline



Fig. 1: MIC compression pipeline. Raw 16-bit pixels are delta-encoded (avg of top and left neighbors), run-length encoded (same/diff runs), and entropy-coded with a 16-bit-native FSE/tANS coder.

Encoding example. For the residual code sequence [5, 5, 5, 5, 9, 10, 11, 11, 11], the RLE output is: same($c = 4$, value=5), diff($c = 2$, values=9, 10), same($c = 3$, value=11). The encoded stream contains five uint16 words: [4, 5, 32769, 9, 10, 3, 11], where $32769 = \text{midCount} + 2$ signals a diff run of length 2.

Why 16-bit RLE matters. A run of 1,000 identical 16-bit residuals $v = 0x0103$ is encoded by MIC as two uint16 words (a count and the value), regardless of v 's byte pattern. A byte-oriented RLE sees the interleaved stream 0x03, 0x01, 0x03, 0x01, ...: no single byte repeats 1,000 times, so it must emit two interleaved runs or fall back to literals.

Same-run fast path. Same-value runs (often >80% of all runs on smooth medical images after delta coding) are fast-pathed in the decoder to return the cached value without touching the compressed-data slice. Each output symbol in a same run requires only a counter decrement.

D. Large-Alphabet FSE Coding

FSE overview. FSE maintains a single integer *state*. To decode one symbol, the decoder indexes a prebuilt decode table using the current state. Each table entry stores the output symbol, the number of bits n to read from the bitstream, and a base value b for the next state. The decoder reads n bits from the bitstream and computes the next state as $b +$ (bits read). This requires only a table lookup, a bit read, and an addition per symbol; no divisions or multiplications are needed. **Encoding** processes symbols in reverse order (last symbol first), building the compressed bitstream from the end; **decoding** reads bits forward.

The RLE output is entropy-coded using table-based ANS (tANS), implemented as Finite State Entropy. MIC extends FSE to support up to 65,535 distinct symbols. As noted in Section I, one 16-bit code value is reserved as the overflow delimiter, leaving at most 65,535 codes available for residual symbols. This contrasts with the 4,096-symbol limit in Zstandard's FSE implementation, which was designed for byte-oriented streams.

Dynamic table sizing. Encode and decode tables are sized to the actual observed symbol range rather than the theoretical maximum. For 8-bit images stored in 16-bit containers

(common in MR), this reduces the working set from 512 KB to approximately 2 KB.

E. Adaptive Table Size Selection

The FSE *tableLog* parameter determines the size of the decode table (2^{tableLog} entries) and the precision of symbol probability estimates. MIC automatically selects *tableLog* based on the number of active symbols and observation density:

- *symbolLen* = number of distinct symbols observed in the RLE output stream (equivalently, the observed symbol range: max value minus min value plus 1, counting all positions in that range even if some go unused).
- *inputLength* = total number of symbols in the RLE output stream fed to the FSE coder.
- *symbolDensity* = $\text{inputLength} / \text{symbolLen}$, i.e., average observations per symbol slot.

Selection rules (applied in order; later rules override earlier ones):

- 1) Default: *tableLog* = 11.
- 2) If *symbolLen* > 128 **and** *symbolDensity* > 32: set *tableLog* = 12.
- 3) If *symbolLen* > 512 **and** *symbolDensity* > 16: set *tableLog* = 13 (overrides rule 2 when both conditions hold).

A larger *tableLog* improves probability precision for the FSE model but increases table memory from $2^{11} \times 4 = 8$ KB to $2^{13} \times 4 = 32$ KB. Rule 2 activates on images with a large active symbol set and enough data per symbol for reliable probability estimates. Rule 3 extends this to images with very broad symbol ranges (e.g., mammography after delta coding), where precision gains outweigh the memory cost even at lower density.

Table IV presents the *tableLog* ablation across selected images.

Nine images benefit from TL=12 or TL=13 (CR +6.3%, MG1 +9.9%, MG2 +9.9%, MR2 +4.6%, RG2 +9.9%, RG3 +7.8%, XA1 +4.4%, MR3 +0.9%, NM1 +1.9%); the adaptive policy correctly identifies all beneficial cases. Decompression throughput is unaffected by *tableLog* choice (<5% variation).

TABLE IV: TableLog ablation: forced TL=11/12/13 vs. adaptive selection (selected images, ARM64 reference platform).

Image	TL=11	TL=12	TL=13	Adaptive
CR	3.47×	3.63×	3.69×	3.69×
MG1	8.00×	8.57×	8.79×	8.79×
MG2	7.98×	8.55×	8.77×	8.77×
MR2	3.14×	3.24×	3.28×	3.28×
RG2	3.85×	4.12×	4.23×	4.23×
RG3	5.64×	5.95×	6.08×	6.08×

IV. MULTI-STATE ENTROPY DECODING

A. The Serial Dependency Problem

A conventional single-state table-based ANS decoder contains a serial dependency chain. Each decoded symbol requires one table lookup followed by one bit read, and the next state depends on both results:

$$\begin{aligned} e &= \text{decTable}[s_n], \\ s_{n+1} &= e.\text{base} + \text{getBits}(e.\text{nbBits}), \end{aligned} \quad (4)$$

where s_n is the current decoder state, `decTable` is the prebuilt decode table, $e.\text{base}$ is the base value for the next state stored in the table entry, $e.\text{nbBits}$ is the number of bits to consume from the bitstream, and `getBits( $k$ )` reads the next k bits from the compressed stream and returns them as an integer. Because s_{n+1} depends on e , which depends on s_n , the computations for symbol $n+1$ cannot begin until symbol n is fully resolved. With a table-lookup latency of approximately $L \approx 4$ cycles on modern out-of-order CPUs (L1 cache hit), the loop is limited to roughly one symbol per L cycles regardless of available execution units. Modern CPUs have 4–6 execution ports; a single-state ANS decoder uses exactly one dependency chain, leaving 75–83% of the hardware capacity idle.

B. Interleaved Decoding

MIC breaks this dependency by running k independent state machines on interleaved symbol positions. With k chains (we evaluate $k \in \{2, 4, 8\}$; the eight-chain configuration is the headline variant), chain j reconstructs the positions $j, j+k, j+2k, \dots$, e.g. for $k=8$:

chain A : 0, 8, 16, ...
 chain B : 1, 9, 17, ...
 ...
 chain H : 7, 15, 23, ...

The chains are completely independent: each has its own state variable, its own sequence of table lookups, and its own bit-extraction operations. The CPU’s out-of-order engine sees k simultaneous table-lookup address streams.

Bitstream format. The encoder partitions the output symbol sequence into k interleaved subsequences, encodes each independently (producing k FSE-compressed streams), and concatenates them. All chains share the same FSE decode table (built from the joint histogram over all symbols). Initial state values for each chain are stored as part of the per-chain header. Output reordering is performed by the decoder, which writes

reconstructed symbols to positions $k \cdot t + j$ from chain j at inner-loop iteration t .

Latency model. Let L be the table-lookup latency (cycles). Single-state FSE processes 1 symbol per L cycles. The k -state decoder processes k symbols per L cycles (amortized):

TABLE V: Multi-state decoder latency model and empirical speedup.

States (k)	Amortized latency	Theoretical	Empirical (ARM64)
1	L cy/sym	1.0×	1.0×
2	$L/2$ cy/sym	2.0×	1.3×
4	$L/4$ cy/sym	4.0×	1.5×
8	$L/8$ cy/sym	8.0×	1.6×

The empirical speedup is below theoretical maximum due to: (1) bit-reader refill overhead shared across all chains, (2) instruction cache pressure from the unrolled loop, and (3) memory bandwidth limits on compressed stream reads [28].

C. Correctness and Signaling

Correctness follows from partitioning the output sequence into independent subsequences during encoding and reversing that assignment during decoding; each chain operates on a strict subset of positions with no cross-chain data dependency.

The bitstream signals the decoding mode with a 2-byte magic header: 0xFF, 0x02 for two-state; 0xFF, 0x04 for four-state; 0xFF, 0x84 for eight-state; followed by a 4-byte symbol count. In the single-state format, the first byte encodes `tableLog` (the value that determines the decode table size as 2^{tableLog} entries).

The `tableLog` is bounded by the symbol alphabet size: for a 16-bit symbol alphabet of at most 65,535 symbols, a table of $2^{16} = 65,536$ entries gives one slot per symbol, which is the practical minimum for meaningful probability quantization. The implementation supports `tableLog` up to 17 (giving 131,072 entries, approximately $2\times$ the alphabet size, for better probability resolution on large alphabets). Since the maximum valid `tableLog` value is 17 and a single byte can represent at most 255, the value 0xFF (= 255 > 17) is not a valid single-state `tableLog` and can safely be reused as the extended-format marker. All three formats coexist without ambiguity.

Platform-specific kernels. On AMD64, Bit Manipulation Instruction Set 2 (BMI2) `SHLQ/SHRQ` instructions provide 3-operand variable shifts without contention on the CL register (the x86 shift count register used by traditional variable shifts); on ARM64, `LSLV/LSRV` (logical shift left/right variable) serve the same role natively. Pure-Go implementations cover all other targets.

D. Multi-State FSE Decoder Implementation

zeroBits flag. The FSE decoder uses a fast inner loop that reads a fixed number of bits per symbol. However, when any symbol’s normalized count exceeds half the table size (i.e., its probability exceeds 50%), some state transitions must emit zero bits—the decoder would call `getBits(0)`. A boolean flag named `zeroBits` (set at table-build time in

the implementation) tracks this condition; Pure-Go code must branch on it to avoid a zero-shift undefined behavior, and the resulting conditional branch breaks the fast path. The AMD64 BMI2 assembly function `fse4StateDecompKernel` (described below) uses the `SHRXQ` instruction, which handles a shift count of zero correctly without branching, so it sidesteps the `zeroBits` check entirely.

`fse4StateDecompKernel` is a hand-written AMD64 assembly routine in the MIC implementation that runs all four interleaved FSE decode chains in a single tight loop. It uses BMI2 `SHLXQ`/`SHRXQ` instructions for register-saving variable shifts (avoiding contention on the CL shift-count register used by legacy x86 variable-shift instructions) and issues four simultaneous table-lookup address computations to keep the CPU’s out-of-order execution engine fully occupied.

V. EXPERIMENTAL METHODOLOGY

A. Dataset

The evaluation uses 21 DICOM images spanning nine modalities, drawn from publicly available datasets: the NEMA Working Group 4 (WG-04) compression sample images [31], the NEMA 1997 CD, and the Clunie Breast Tomosynthesis Case 1 (Table VI).

TABLE VI: Test dataset: 21 clinical DICOM images spanning nine modalities.

Image	Modality	Dimensions	Raw (MB)
MR	Brain MRI	256×256	0.13
CT	CT scan	512×512	0.50
CR	Comp. radiography	2140×1760	7.18
XR	X-ray	2048×2577	10.1
MG1	Mammography	2457×1996	9.35
MG2	Mammography	2457×1996	9.35
MG3	Mammography	4774×3064	27.3
MG4	Mammography	4096×3328	26.0
CT1	CT scan	512×512	0.50
CT2	CT scan	512×512	0.50
MG-N	Mammography	3064×4664	27.3
MR1	Brain MRI	512×512	0.50
MR2	Brain MRI	1024×1024	2.00
MR3	Brain MRI	512×512	0.50
MR4	Brain MRI	512×512	0.50
NM1	Nuclear med.	256×1024	0.50
RG1	Radiography	1841×1955	6.86
RG2	Radiography	1760×2140	7.18
RG3	Radiography	1760×1760	5.91
SC1	Sec. capture	2048×2487	9.71
XA1	Fluoroscopy	1024×1024	2.00

This dataset spans multiple modalities and image sizes but remains modest for strong generalization claims.

B. Baselines

The current study compares MIC with:

- **HTJ2K** (OpenJPH [29]),
- **JPEG-LS** (CharLS [30]), and
- **Delta + Zstandard** (level 19).

All three baselines are called in-process (no subprocess or file-I/O overhead) using Go’s C foreign-function interface (CGO) for the C-library codecs.

C. Metrics

The paper reports:

- **Compression ratio:** uncompressed size / compressed size.
- **Decompression throughput:** uncompressed pixel bytes reconstructed per second (MB/s).
- **Encoding throughput:** uncompressed pixel bytes processed per second (MB/s).

D. Benchmark Procedure

All codecs were measured using in-process library calls on the same machine to avoid file-I/O and subprocess overhead. The Go testing framework’s `-benchtime=10x` flag was used (10 iterations per benchmark; each iteration processes the full image). Run-to-run variance is <2% on ARM64 and <5% on AMD64.

Reference platforms. The paper reports two architectures, ARM64 and AMD64. All ARM64 and AMD64 numbers in this paper come from the reference machines listed below; this is the only place where specific CPU models appear. Captions, prose, and tables elsewhere refer only to the ARM64/AMD64 categories.

- **ARM64 reference** (all C and Go benchmarks): Apple M4 Pro (14-core: 10 performance cores plus 4 efficiency cores), macOS. Go compiler (gc) with default optimization; CGO-linked C components compiled with `-O3`.
- **AMD64 reference** (all C and Go benchmarks): AWS EC2 `c8i.4xlarge` (Intel Xeon 6 / Granite Rapids, 16 vCPU), Ubuntu Linux 24.04. Go compiler (gc) with default optimization; C components compiled with `-O3 -march=native`.
- **ARM64 JavaScript reference** (Tables XII and XIII only): Apple M2 Max (12-core ARM64) workstation, macOS, Node.js v24.8.

Codecs evaluated. The headline single-threaded tables in Section VI compare three codecs:

- **MIC** — the proposed codec. In every headline table the “MIC” column refers to the fastest available C decoder on that platform: on ARM64 it is the eight-state C decoder, and on AMD64 it is the same eight-state C decoder with AVX-512 acceleration of the RLE expansion and delta inverse stages that follow FSE (Section IV). Both implementations read the same compressed bytes; the only platform-specific difference is the SIMD width used for the RLE and delta passes.
- **HTJ2K** — High-Throughput JPEG 2000 via the OpenJPH reference implementation [29], compiled with the same C optimisation flags as MIC. We note that the OpenJPH source tree ships vectorised HT-block decoders for x86 only (SSSE3, AVX2, AVX-512) and contains no NEON or SVE kernel paths; on ARM64 the runtime dispatcher in OpenJPH’s codestream layer falls back to the scalar block decoder. Reported ARM64 HTJ2K throughput should therefore be read as the performance of the OpenJPH reference rather than of HTJ2K’s architectural ceiling on ARM.

- **JPEG-LS** — the standardised predictive lossless codec via CharLS [30]. Reported on both reference platforms.

MIC has several additional implementation variants (pure Go, single-state, scalar C, wavelet-based) which are useful for understanding the design trade-offs but are not central to the codec-to-codec comparison. They are summarised separately in Section VI-E. A strip-parallel mode that decodes one image using multiple threads is also implemented and reported separately in Section VI-G for completeness.

E. JavaScript and Browser Evaluation

A pure JavaScript decoder (~20 KB ES module, zero Node Package Manager (npm) dependencies) and a Go WebAssembly build (~2.5 MB) were implemented. JavaScript performance measurements reported in Section VI were obtained on the ARM64 JavaScript reference platform listed in Section V-D; the C and Go benchmarks use the ARM64 and AMD64 references in the same section.

VI. RESULTS

A. Compression Ratio

Table VII presents lossless compression ratios for all codec variants across the 21-image dataset.

Analysis. JPEG-LS achieves the highest lossless ratio on all 21 images, as expected given its context-adaptive MED predictor. MIC’s design target is not ratio alone but the ratio-throughput-simplicity tradeoff: MIC achieves approximately 91% of JPEG-LS’s geometric mean ratio ($3.12\times$ vs. $3.44\times$) while decoding faster and requiring approximately 1/5 the source lines. Mammography achieves the highest MIC ratios (up to $8.79\times$) due to smooth tissue regions producing near-zero RLE runs.

PICS strip adaptation. Each PICS strip receives a dedicated FSE table that specializes to its local residual distribution. On large images, this per-strip specialization improves ratio slightly. Formally, if S_k denotes the RLE symbol subsequence for strip k , $|S_k|$ its length, and $H(S_k)$ its Shannon entropy (bits per symbol), then

$$\sum_k |S_k| \cdot H(S_k) \leq |S| \cdot H(S), \quad (5)$$

where S is the full sequence and the inequality is strict when residual statistics differ across strips. In plain terms: a table that is specialized to a strip’s own pixel-value distribution is at least as efficient as a single global table, and strictly better when different image regions have different statistical properties.

B. Encoding Throughput

Summary. MIC encodes faster than HTJ2K, JPEG-LS, and Δ +Zstandard on the large majority of the 21 medical images we tested on both reference platforms. On ARM64, MIC leads on every image, running at roughly 1.9 – $3.2\times$ HTJ2K’s throughput, 3 – $7\times$ JPEG-LS’s throughput, and 35 – $240\times$ Δ +Zstd-19’s throughput. On AMD64, MIC leads on 18 of 21 images at roughly 1.0 – $1.6\times$ HTJ2K and 3 – $6\times$ JPEG-LS; HTJ2K narrowly wins on the two largest mammography slabs

(MG1, MG2) and on RG3. The detailed per-image numbers are in Table VIII.

The throughput advantage is primarily structural: MIC’s encoder is a single pass over the pixels (a spatial predictor, a run-length pass, and a table-driven entropy coder), whereas HTJ2K’s block-based EBCOT and JPEG-LS’s context-adaptive Golomb-Rice coder each do more work per pixel. Δ +Zstd-19 sits at the opposite end of the spectrum: the maximum-ratio Zstandard preset runs an aggressive long-range LZ77 search which dominates encode time, yielding the 40 – $200\times$ throughput gap at the cost of consistently worse compression ratio than MIC across all 21 images. The fact that the gap to HTJ2K narrows on AMD64 reflects the AMD64 baseline codecs being themselves well-tuned with vector instructions, so the relative speedup MIC obtains from its simpler pipeline is smaller.

C. Decompression Throughput—ARM64

Summary. On ARM64, MIC is the fastest single-threaded decoder on 20 of the 21 medical images we tested; Δ +Zstandard-19 narrowly leads only on CT (472 vs. 466 MB/s, $+1.3\%$). MIC decodes roughly 1.5 – $2.6\times$ faster than HTJ2K, 1.1 – $1.7\times$ faster than Δ +Zstd-19, and 2.0 – $4.5\times$ faster than JPEG-LS, with the largest margins on the smaller-image modalities. The per-image numbers are in Table IX.

The advantage is again structural. MIC’s decoder is dominated by a table-driven entropy stage: one table lookup, one bit read, and one addition per output symbol, with no data-dependent branches that the processor’s branch predictor cannot anticipate. HTJ2K’s block coder and JPEG-LS’s context model each introduce data-dependent branches whose mispredictions a wide out-of-order engine can only partially hide. Δ +Zstandard is the closest architectural sibling — it shares MIC’s spatial predictor and uses FSE for entropy coding — but operates on bytes rather than 16-bit residuals, so it pays an extra factor-of-two per-symbol coding cost on high-bit-depth data while matching MIC’s branch-light decode shape. The combination of better ratio (5 – 22% over Δ +Zstd on every image) and faster decode (20/21 images) is the empirical signature of the 16-bit-native design. JPEG-LS retains slightly better compression ratios (around 10% better on average), so the trade-off when choosing between MIC and JPEG-LS is between decode latency and archival storage cost.

D. Decompression Throughput—AMD64

Summary. On AMD64, MIC is the fastest single-threaded decoder on all 21 of the 21 medical images we tested. The margin over the runner-up ranges from $1.01\times$ on MG1 (the closest contest, 913 vs. 907 MB/s) to $1.62\times$ on XR. MIC decodes roughly 1.0 – $1.6\times$ HTJ2K, 1.1 – $1.6\times$ Δ +Zstd-19, and 3.9 – $6.9\times$ JPEG-LS. The per-image numbers are in Table IX.

This is a stronger result than the four-state SIMD variant produced on this platform: switching the AMD64 “MIC” column from four-state to eight-state turned three of the previously narrow HTJ2K wins (MG1, MG2, MG4) and two narrow Δ +Zstandard wins (CT, CT2) into MIC wins. The win has two sources. First, the eight-state inner loop exposes eight

TABLE VII: Lossless compression ratios: all codec variants, 21 images. Bold = best ratio per row.

Image	Raw (MB)	Δ +Zstd-19	MIC	Wavelet	PICS-4	PICS-8	HTJ2K	JPEG-LS
MR	0.13	1.95×	2.35×	2.38×	2.28×	2.21×	2.38×	2.52×
CT	0.50	2.05×	2.24×	1.67×	2.15×	1.96×	1.77×	2.68×
CR	7.18	3.24×	3.69×	3.81×	3.70×	3.71×	3.77×	3.96×
XR	10.1	1.43×	1.74×	1.76×	1.75×	1.76×	1.67×	1.76×
MG1	9.35	7.20×	8.79×	8.67×	8.84×	8.87×	8.25×	8.91×
MG2	9.35	7.21×	8.77×	8.65×	8.83×	8.85×	8.24×	8.90×
MG3	27.3	2.04×	2.24×	2.32×	2.31×	2.34×	2.22×	2.38×
MG4	26.0	3.10×	3.47×	3.59×	3.59×	3.62×	3.51×	3.71×
CT1	0.50	2.47×	2.79×	2.49×	2.54×	2.29×	2.70×	3.19×
CT2	0.50	3.24×	3.49×	2.87×	3.11×	2.72×	3.29×	4.54×
MG-N	27.3	2.04×	2.24×	2.32×	2.31×	2.34×	2.23×	2.38×
MR1	0.50	1.79×	2.09×	2.14×	2.10×	2.08×	2.13×	2.30×
MR2	2.00	2.77×	3.28×	3.34×	3.31×	3.31×	3.35×	3.52×
MR3	0.50	3.47×	3.93×	4.09×	3.89×	3.84×	4.33×	4.51×
MR4	0.50	3.58×	4.12×	4.18×	4.09×	4.03×	4.21×	4.49×
NM1	0.50	4.88×	5.15×	5.02×	5.26×	5.28×	5.76×	6.28×
RG1	6.86	1.45×	1.70×	1.70×	1.70×	1.69×	1.63×	1.72×
RG2	7.18	3.68×	4.23×	4.32×	4.28×	4.30×	4.32×	4.51×
RG3	5.91	5.46×	6.08×	6.82×	6.11×	6.12×	6.99×	7.31×
SC1	9.71	3.46×	3.71×	3.70×	3.73×	3.74×	3.85×	4.73×
XA1	2.00	4.37×	5.01×	4.94×	5.04×	5.03×	4.88×	5.39×

TABLE VIII: Single-threaded encoding throughput (MB/s). Codec baselines as defined in Section V; MIC is the eight-state C encoder. Bold marks the fastest codec per row on each platform.

Image	ARM64				AMD64			
	MIC	Δ +Zstd	HTJ2K	JPEG-LS	MIC	Δ +Zstd	HTJ2K	JPEG-LS
MR	504	8	273	97	338	5	229	62
CT	529	8	200	139	292	5	206	97
CR	569	3	204	115	322	2	284	75
XR	690	3	248	125	405	3	250	84
MG1	1,168	14	456	315	546	8	647	209
MG2	1,172	14	454	315	541	8	646	209
MG3	706	3	234	143	332	2	251	104
MG4	993	7	373	215	431	6	415	150
CT1	576	11	238	178	329	7	277	125
CT2	561	8	207	172	286	5	266	125
MG-N	718	2	235	143	339	2	243	104
MR1	677	11	215	120	344	6	244	86
MR2	737	9	231	131	397	4	300	84
MR3	901	22	299	187	428	12	353	123
MR4	711	8	228	175	358	6	318	141
NM1	663	7	247	170	336	5	312	116
RG1	659	5	246	95	392	4	250	65
RG2	811	6	271	151	416	3	356	99
RG3	762	6	285	189	397	4	407	126
SC1	765	6	256	223	407	4	309	168
XA1	693	7	232	152	358	4	310	97

independent FSE state-update chains, doubling the work the wide-issue Granite Rapids core can keep in flight per cycle without any SIMD assistance. Second, the AVX-512 widening of the RLE-expand and delta-inverse passes (32 uint16s per store, eight times wider than scalar) shifts a meaningful fraction of post-FSE time from load-store-bound serial code to vector throughput. The same code path runs on ARM64 and exhibits the same qualitative improvement (Table IX, ARM64 columns); the larger AMD64 absolute gain reflects the AVX-512 lane width relative to ARM64 NEON.

E. Implementation Variants of MIC

Summary. The headline tables use a single MIC column for clarity: the eight-state C decoder on ARM64 and the same decoder with AVX-512 acceleration of the RLE expansion and delta inverse passes on AMD64. MIC has several other implementation variants that exist for different deployment situations rather than for headline competition. They are described briefly here so the design space is on record.

- **Pure Go reference.** The full encoder and decoder run in pure Go through the standard gc compiler with no special build flags. It decompresses at approximately 40–60% of the C variant’s throughput on both reference platforms (e.g. 248 MB/s on the MR test image on ARM64 vs. 564 MB/s for the C decoder). The Go reference is useful where the deployment cannot accept a CGO dependency (mobile, WebAssembly via TinyGo, or strictly-sandboxed runtimes).
- **Single-state, two-state, and four-state decoders.** Variants of the same C decoder with $k \in \{1, 2, 4\}$ entropy-coder states instead of eight. Each step down from eight to four to two to one progressively slows the FSE stage as instruction-level parallelism is removed (Table X quantifies the FSE stage in isolation across the four configurations). These variants exist as baselines for the multi-state ablation in Section IV.
- **Wavelet-based MIC.** An alternative MIC pipeline that replaces the spatial predictor with a five-level reversible

TABLE IX: Single-threaded decompression throughput (MB/s). Codec baselines as defined in Section V; MIC is the eight-state C decoder (AVX-512 acceleration of the RLE and delta inverse stages on AMD64). Bold marks the fastest codec per row on each platform.

Image	ARM64				AMD64			
	MIC	Δ +Zstd	HTJ2K	JPEG-LS	MIC	Δ +Zstd	HTJ2K	JPEG-LS
MR	564	331	301	146	552	357	439	91
CT	466	472	356	187	453	384	329	111
CR	692	564	367	206	566	475	439	123
XR	655	572	366	150	617	508	386	88
MG1	829	691	681	547	913	599	907	332
MG2	873	684	686	557	920	599	896	333
MG3	678	473	362	201	556	358	390	117
MG4	815	581	533	259	676	466	611	157
CT1	571	502	373	238	525	425	390	143
CT2	648	575	372	222	512	456	383	133
MG-N	657	469	360	201	545	355	391	117
MR1	681	437	354	154	603	384	397	94
MR2	747	505	390	224	644	464	491	132
MR3	859	507	453	329	773	485	543	192
MR4	775	523	304	230	632	480	446	153
NM1	866	595	404	277	598	486	467	172
RG1	561	413	354	126	550	347	363	80
RG2	755	574	408	255	730	497	535	154
RG3	740	655	469	320	723	552	628	196
SC1	696	587	384	280	690	484	456	183
XA1	759	611	388	278	688	534	477	165

5/3 integer wavelet transform and feeds the wavelet coefficients into the same run-length and entropy stages. On AMD64 the predict/update lifting steps use AVX2 kernels; on ARM64 the kernels are scalar. The wavelet variant is competitive on AMD64 for a small number of images (notably CT2 and RG1) but is consistently slower than the predictor-based MIC on the medical-imaging corpus, because the wavelet transform pays a cost that the spatially smooth medical images do not reward. It is included in our test suite as a back-end option but does not appear in the headline tables.

The eight-state C decoder is therefore both the simplest and the fastest configuration on the evaluated corpus. The other variants document the implementation space rather than competing for the fastest column.

F. FSE Stage Microbenchmarks

Table X reports FSE decompression throughput measured in isolation (MB/s over the uncompressed RLE symbol stream), decoupling the entropy-coding stage from delta and RLE overhead, for a representative subset of modalities on both platforms. The 1-, 4- and 8-state variants are all pure Go: the only difference is how many independent FSE state chains the decoder interleaves per inner-loop iteration.

The 4-state decoder delivers a geometric-mean speedup of +46% on ARM64 and +8% on AMD64 over the 1-state baseline; the 8-state decoder adds a further +7% on ARM64 and +2% on AMD64 over 4-state. The ARM64 trajectory is roughly monotonic with state count: each doubling of independent chains continues to unlock more out-of-order parallelism. The AMD64 trajectory flattens earlier; the Granite Rapids core already extracts substantial ILP from the 1-state path through speculative scheduling, so the additional chains help less. Note that the FSE microbench isolates the entropy stage from RLE and delta overheads — in the full pipeline (Table IX), AMD64 still gains +9–10% from 4-state to 8-state

TABLE X: FSE decompression throughput (MB/s) for representative modalities, single-thread pure-Go decoder. Increasing the number of independent state chains exposes more instruction-level parallelism to the out-of-order engine.

Image	ARM64 (MB/s)			AMD64 (MB/s)		
	1-st	4-st	8-st	1-st	4-st	8-st
MR	612	953	765	636	899	759
CT	776	1,872	1,729	974	986	953
CR	3,051	5,749	5,521	1,426	1,581	1,638
XR	3,657	6,072	5,807	1,939	1,890	2,248
MG1	3,604	4,557	5,887	2,075	1,521	1,753
MG-N	3,588	6,378	6,307	1,714	2,294	2,473
NM1	1,798	2,535	2,740	1,091	1,274	1,522
RG1	1,786	3,492	5,412	1,198	1,717	2,104
Geomean (21)	2,349	3,422	3,667	1,464	1,575	1,614

because the 8-state SIMD variant additionally widens the RLE-expand and delta-inverse passes to AVX-512.

G. Multi-threaded Decoding

Summary. MIC supports an optional strip-parallel decode mode in which the image is split into horizontal strips, each strip is decoded independently in a separate thread, and the rows are concatenated. We report it here for completeness rather than as a head-to-head comparison: HTJ2K and JPEG-LS are single-threaded in our pipeline, so an N -thread-vs-1-thread comparison would not be meaningful. The take-away is that the strip-parallel decoder scales close to linearly up to four threads on both reference platforms and continues to scale to eight threads on images above approximately 5 MB, reaching 4–4.6 GB/s on the largest mammography images on ARM64. Small images do not benefit: synchronization and per-strip header overhead exceed the per-strip work below roughly half a megabyte.

[†]The MR image (130 KB raw) is too small to benefit from 8-way parallelism: synchronization overhead exceeds the per-strip work, so the 4-strip configuration is faster than 8-strip on both platforms.

TABLE XI: Multi-threaded decompression throughput (MB/s) for MIC’s strip-parallel mode. “MIC” is the single-threaded baseline from Table IX reproduced for reference; N -strip columns decode the same compressed bytes in N parallel threads. The MR image is too small for 8-way parallelism on both platforms.

Image	ARM64				AMD64			
	MIC	2-strip	4-strip	8-strip	MIC	2-strip	4-strip	8-strip
MR	702	663	894	702 [†]	481	329	367	294 [†]
CT	470	624	1,157	1,446	368	398	633	714
CR	644	1,052	1,976	3,605	530	758	1,320	2,169
XR	682	1,087	2,024	3,794	530	793	1,515	2,579
MG1	848	1,529	2,529	4,443	889	1,317	1,951	2,890
MG2	877	1,417	2,506	4,585	869	1,356	1,971	3,026
MG3	681	1,112	1,911	3,820	509	838	1,449	2,767
MG4	797	1,326	2,345	4,348	644	1,207	1,897	3,126
CT1	531	819	1,247	1,525	478	458	703	725
CT2	505	864	1,260	1,473	450	481	734	719
MG-N	661	1,147	2,027	3,930	522	913	1,537	2,819
MR1	724	977	1,460	1,851	536	504	761	764
MR2	741	1,137	1,911	3,103	613	778	1,171	1,534
MR3	846	1,205	1,693	2,089	695	591	788	835
MR4	770	997	1,469	1,757	588	540	796	804
NM1	852	1,147	1,682	2,133	586	523	824	762
RG1	476	677	1,203	2,150	477	603	1,101	1,890
RG2	758	1,229	2,156	3,816	661	916	1,473	2,268
RG3	772	1,335	2,343	4,168	672	971	1,618	2,730
SC1	758	1,311	2,206	3,971	627	1,024	1,673	2,541
XA1	759	1,214	1,881	3,191	617	792	1,246	1,748

The scaling characteristics are intuitive. On large images each strip gets enough rows that the per-strip overhead (synchronization plus reading a small strip header) is amortised, and total decode time falls roughly as $1/N$ up to $N = 4$ and continues to fall to $N = 8$. On small images the overhead dominates and adding threads slows things down. As a practical matter, applications that decode many small images concurrently should use the single-threaded MIC decoder per request (which is already fast); applications that decode one large image at a time on a multi-core machine should use the strip-parallel mode and pick N between four and eight depending on image size and core count.

H. JavaScript Runtime Results

Table XII reports Node.js single-threaded performance; Table XIII reports parallel PICS decoding via `worker_threads`.

TABLE XII: JavaScript decoder throughput (4-state) on the ARM64 JavaScript reference (Section V-D). Median over 20 iterations.

Image	Dimensions	Ratio	ms	MB/s
MR	256×256	2.35×	3.0	42
CT	512×512	2.24×	13.5	37
CR	2140×1760	3.69×	161.2	45
MG1	2457×1996	8.79×	80.9	116
MG3	4774×3064	2.29×	638.4	44

MG1 (9.35 MB) decompresses in 19.4 ms at 483 MB/s using 12 workers. A web-based DICOM viewer can fetch a `.mic` file directly from object storage and decode it client-side without server-side compute or transcoding latency.

Limitation. These measurements were obtained in Node.js `worker_threads`, which does not fully replicate browser Web Worker performance. Full browser performance benchmarking remains future work.

TABLE XIII: PICS parallel JavaScript decoder (`worker_threads`) on the ARM64 JavaScript reference. Speedup relative to 1 worker.

Image	Strips	ms	MB/s	Speedup
MR	4	1.3	94	2.57×
CT	4	5.0	101	2.95×
CR	8	31.7	227	5.53×
MG1	8	19.4	483	4.36×

VII. CONCLUSION

Bottom line. Across 21 clinical DICOM images spanning nine modalities, MIC is the fastest single-threaded lossless decoder on every ARM64 image and on every AMD64 image we tested, while compressing within 1% of the geometric-mean ratio of High-Throughput JPEG 2000. On the fastest-encoding side it is even more decisive: MIC encodes faster than both HTJ2K and JPEG-LS on every image on both platforms. The codec is implemented in roughly 1,100 lines of Go, about 18 times fewer source lines than the OpenJPH reference implementation of HTJ2K.

Reading the results. Three observations tie the per-table numbers together. First, the gain that MIC obtains comes primarily from its design choice to operate directly on 16-bit residuals: the entropy coder sees the 10- to 16-bit symbols produced by the spatial predictor as a single alphabet rather than splitting each pixel into two bytes. This sidesteps the structural cost of byte-splitting (5–22% compared to a Delta+Zstandard baseline on the same predictor) and at the same time removes the per-symbol branches that the byte-splitting codecs spend cycles on during decode. Second, on the very largest mammography images on AMD64, HTJ2K’s block coder recovers the throughput lead by amortising its per-block setup over many pixels; MIC remains within 0–20% of HTJ2K’s throughput on those images. Third, multi-state interleaved entropy decoding gives MIC a further ~45% on

ARM64 and $\sim 43\%$ on AMD64 without any change to the compressed bytes (Table X); this is the gain that makes MIC’s C decoder competitive with HTJ2K’s vectorised block coder.

Deployment guidance. For applications that decode one large image at a time on a multi-core machine, MIC’s strip-parallel mode delivers the highest absolute decode rate (Section VI-G). For applications that decode many small images concurrently — thumbnail rendering, multi-frame breast tomosynthesis playback, real-time PACS panning — single-threaded MIC is the right default: each request gets approximately 700–900 MB/s on ARM64 and 800–1,200 MB/s on AMD64 without contending for threads with other requests. HTJ2K remains a sensible choice where compatibility with existing DICOM transfer syntaxes is a hard requirement, and is specifically competitive on large mammography images on AMD64. JPEG-LS is the right choice where archival storage cost dominates the deployment budget — it retains a small ratio advantage (about 10% geometric mean) but is 1.6–5.7 $\times$ slower to decode on ARM64. For deployment situations that cannot accept a native CGO dependency, the pure-Go reference and the 20 KB JavaScript decoder offer implementations of the same compressed format at roughly half the throughput of the C variant.

A. Codec Selection Framework

The deployment guidance above is qualitative. We close with a more structured comparison that combines the quantitative metrics of Section VI with the platform-availability dimensions a deployment team weighs in practice. The three scenarios used below (low-latency viewing, archival storage, and browser delivery) follow the deployment taxonomy laid out for DICOM compression by Clunie [38] and reviewed for current medical-imaging practice by Liu *et al.* [24]. The criterion set — compression ratio, encode and decode throughput, AMD64 and ARM64 native availability, browser compatibility, and deployment ease — generalises the four classical evaluation axes of Sayood [39] (ratio, complexity, memory, speed) to include the platform-coverage dimensions that medical-imaging products contend with in practice. Source-line counts and the choice of implementation language are deliberately excluded from the criterion set so that the matrix captures what an integrator sees externally, not how the codec happens to be built internally. Two complementary tools are used: (i) a *Pareto-frontier* view, the dominant method for lossless codec comparison in the wider compression literature [33]–[35], which surfaces the two-axis ratio–throughput trade-off without committing to fixed weights; and (ii) a *weighted decision matrix* that extends the analysis to the platform and deployment criteria above, following the multi-criteria scoring approach common in software- and codec-selection studies [36], [37].

Pareto frontier. Fig. 2 plots compression ratio against single-threaded ARM64 decompression throughput for the four codecs benchmarked in Table IX. Each codec contributes 21 small points (per-image) and one large marker (geometric mean). A codec is Pareto-optimal when no other codec achieves both higher ratio *and* higher throughput. Two codecs




Fig. 2: Ratio vs. decompression throughput on ARM64. Small dots are per-image; large markers are geometric means. MIC and JPEG-LS occupy the Pareto frontier; HTJ2K and Δ +Zstd-19 are dominated on both axes.

sit on the frontier: MIC, which delivers the highest geometric-mean decode throughput (700 MB/s) at a ratio within 1% of HTJ2K’s, and JPEG-LS at the opposite end with the highest ratio (3.82 $\times$) but the lowest throughput (238 MB/s). HTJ2K and Δ +Zstd-19 are dominated by MIC on both axes — strictly inside the frontier — so they are preferable only when criteria not captured in the two-axis view (DICOM transfer-syntax compatibility for HTJ2K, general-purpose ubiquity for Zstd) outweigh raw ratio and throughput.

Use-case decision matrix. Pareto analysis is limited to two quantitative axes. Real deployment choices also depend on which CPU architectures and runtime environments a codec covers, and on the friction of integrating it into a product build. Table XIV normalises seven criteria so that the best codec on each row scores 1.0; weighted sums for three representative scenarios are reported in the final three rows. The quantitative rows are derived from the geometric means in Tables VII, VIII, and IX; the platform and deployment rows use the rubric in the table caption.

Interpretation. All three codecs reach parity on native AMD64 and ARM64 availability; those rows confirm the baseline but do not differentiate the codecs. Differentiation comes from the four quantitative rows plus browser support and deployment friction. MIC achieves the highest weighted score in all three scenarios under these neutral weights: 0.98 versus 0.75/0.70 for the viewer, 0.96 versus 0.78/0.80 for archival, and 0.99 versus 0.78/0.74 for browser delivery. The gap narrows on archival, where the ratio weight rises to 0.50 and JPEG-LS edges past HTJ2K (0.80 vs. 0.78) on the strength of its top compression ratio; JPEG-LS still does not overtake MIC because its decode throughput remains roughly 3 $\times$ slower even when ratio is the dominant criterion. HTJ2K is never the highest-scoring choice under these weights, but remains the default whenever DICOM transfer-syntax compatibility is a hard constraint not modelled in the matrix; in such cases the matrix serves to quantify the cost of that compatibility requirement rather than to recommend a different codec. Practitioners with different priorities can re-evaluate Table XIV by adjusting w_P ,

TABLE XIV: Codec selection matrix. Each criterion is normalised so the best codec on that row scores 1.00. Quantitative rows are geometric means across the 21-image set. AMD64 / ARM64 native rows score 1.00 when a production-grade native binary is available for that platform. Browser compatibility: shipped purpose-built JS or WASM decoder used in production = 1.00; WASM port of the native library available and used in clinical viewers (OpenJPH-JS, charls-js via Cornerstone) = 0.85; no browser path = 0.0. Deployment ease: zero external runtime dependencies = 1.00; ubiquitous system dependency (e.g. libzstd) = 0.90; C/C++ library requiring CMake build and link integration = 0.70. Weight columns w_P , w_A , w_B reflect typical priorities for (P) PACS viewer / low-latency decode, (A) archival storage, (B) browser delivery.

Criterion	w_P	w_A	w_B	MIC	HTJ2K	JPEG-LS
Compression ratio	0.20	0.50	0.10	0.91	0.90	1.00
Decompression throughput (ARM64)	0.30	0.10	0.15	1.00	0.58	0.34
Encoding throughput (ARM64)	0.05	0.10	0.05	1.00	0.36	0.22
AMD64 native	0.10	0.05	0.05	1.00	1.00	1.00
ARM64 native	0.10	0.05	0.05	1.00	1.00	1.00
Browser compatibility	0.05	0.00	0.40	1.00	0.85	0.85
Deployment ease	0.20	0.20	0.20	1.00	0.70	0.70
Weighted score (P) viewer	—	—	—	0.98	0.75	0.70
Weighted score (A) archival	—	—	—	0.96	0.78	0.80
Weighted score (B) browser	—	—	—	0.99	0.78	0.74

w_A , or w_B ; the per-row normalised scores remain valid across any weighting scheme.

Future work should strengthen the evidence base with larger datasets, shuffle-based general-purpose baselines, formal statistical reporting, and direct browser benchmarking. Several specific items: (i) measure JPEG-LS, Δ +Zstd-19, and PICS-C-2 throughput on AMD64 so the AMD64 encoding (Table VIII) and decoding (Table IX) tables are column-for-column comparable to the ARM64 tables; (ii) add a NEON predict/update kernel for the wavelet variant on ARM64; and (iii) extend the JavaScript decoder evaluation to the AMD64 platform.

MIC is open-source and available at <https://github.com/pappuks/medical-image-codec>.

ACKNOWLEDGMENT

The author thanks the NEMA Working Group 4 for making the DICOM compression sample images publicly available, and the developers of OpenJPH and CharLS for their open-source implementations used in the comparative evaluation.

CONFLICT OF INTEREST

The author declares no competing financial interests or personal relationships that could have influenced the work reported in this paper.

DATA AVAILABILITY

The test images are drawn from publicly available DICOM datasets: NEMA WG-04 compression samples [31], NEMA 1997 CD, and the Clunie Breast Tomosynthesis Case 1 [32]. The MIC implementation, benchmark harness, and reproduction instructions are available at <https://github.com/pappuks/medical-image-codec>.

REFERENCES

- [1] NEMA, “Digital Imaging and Communications in Medicine (DICOM),” PS3.5—Data Structures and Encoding, PS3.15—Security and System Management Profiles, <https://www.dicomstandard.org/>, 2024.
- [2] D. S. Taubman and M. W. Marcellin, *JPEG2000: Image Compression Fundamentals, Standards and Practice*. Springer, 2002.
- [3] D. S. Taubman, “High throughput JPEG 2000 (HTJ2K): Algorithm, performance evaluation, and potential,” *Proc. SPIE*, vol. 11137, 2019.
- [4] M. J. Weinberger, G. Seroussi, and G. Sapiro, “The LOCO-I lossless image compression algorithm: Principles and standardization into JPEG-LS,” *IEEE Trans. Image Process.*, vol. 9, no. 8, pp. 1309–1324, 2000.
- [5] ISO/IEC 14495-2:2003, “Information technology—Lossless and near-lossless compression of continuous-tone still images: Extensions—Part 2,” 2003.
- [6] J. Alakuijala *et al.*, “JPEG XL next-generation image compression architecture and coding tools,” *Proc. SPIE* 11137, Sep. 2019.
- [7] J. Duda, “Asymmetric numeral systems,” arXiv:0902.0271, 2009.
- [8] J. Duda, “Asymmetric numeral systems: entropy coding combining speed of Huffman coding with compression rate of arithmetic coding,” arXiv:1311.2540, 2013.
- [9] Y. Collet, “Finite State Entropy—A new breed of entropy coder,” <https://github.com/Cyan4973/FiniteStateEntropy>, 2013.
- [10] Y. Collet and M. Kuchera, “Zstandard Compression and the ‘application/zstd’ Media Type,” RFC 8878, IETF, Feb. 2021.
- [11] S. Mahapatra and K. Singh, “An FPGA-based implementation of multi-alphabet arithmetic coding,” *IEEE Trans. Circuits Syst. I*, vol. 54, no. 8, pp. 1678–1686, 2007.
- [12] D. Kosolobov, “Efficiency of ANS Entropy Encoders,” arXiv:2201.02514, 2022.
- [13] R. Bamler, “Understanding Entropy Coding With Asymmetric Numeral Systems (ANS): a Statistician’s Perspective,” arXiv:2201.01741, 2022.
- [14] F. Giesen, “Interleaved entropy coders,” arXiv:1402.3392, 2014.
- [15] F. Giesen, “ryg_rans: Public domain rANS encoder/decoder,” https://github.com/rygorous/ryg_rans, 2014.
- [16] F. Lin, K. Arunrungsirilert, H. Sun, and J. Katto, “Recoil: Parallel rANS Decoding with Decoder-Adaptive Scalability,” *Proc. 52nd ICPP* ’23, pp. 31–40, 2023.
- [17] A. Weissenberger and B. Schmidt, “Massively Parallel ANS Decoding on GPUs,” *Proc. 48th ICPP* ’19, Article 100, 2019.
- [18] X. Wu and N. Memon, “Context-based, adaptive, lossless image coding,” *IEEE Trans. Commun.*, vol. 45, no. 4, pp. 437–444, 1997.
- [19] J. Sneyers and P. Wuille, “FLIF: Free Lossless Image Format based on MANIAC Compression,” *Proc. IEEE ICIP*, 2016.
- [20] D. Le Gall and A. Tabatabai, “Sub-band coding of digital images using symmetric short kernel filters and arithmetic coding techniques,” *Proc. IEEE ICASSP*, pp. 761–764, 1988.
- [21] W. Sweldens, “The Lifting Scheme: A custom-design construction of biorthogonal wavelets,” *Appl. Comput. Harmon. Anal.*, vol. 3, no. 2, pp. 186–200, 1996.
- [22] I. Daubechies and W. Sweldens, “Factoring wavelet transforms into lifting steps,” *J. Fourier Anal. Appl.*, vol. 4, no. 3, pp. 247–269, 1998.
- [23] A. R. Calderbank, I. Daubechies, W. Sweldens, and B.-L. Yeo, “Wavelet transforms that map integers to integers,” *Appl. Comput. Harmon. Anal.*, vol. 5, no. 3, pp. 332–369, 1998.
- [24] F. Liu *et al.*, “The Current Role of Image Compression Standards in Medical Imaging,” *Information*, vol. 8, no. 4, p. 131, Oct. 2017.
- [25] D. A. Clunie, “Lossless Compression of Grayscale Medical Images: Effectiveness of Traditional and State-of-the-Art Approaches,” *Proc. SPIE Medical Imaging*, 2000.
- [26] W3C, “Portable Network Graphics (PNG) Specification (Second Edition),” ISO/IEC 15948:2003, Nov. 2003.
- [27] D. Taubman, A. Naman, M. Smith, P.-A. Lemieux, H. Saadat, O. Watanabe, and R. Mathew, “High Throughput JPEG 2000 for Video Content Production and Delivery Over IP Networks,” *Front. Signal Process.*, vol. 2, Article 885644, Apr. 2022.
- [28] S. Williams, A. Waterman, and D. Patterson, “Roofline: An insightful visual performance model for multicore architectures,” *Commun. ACM*, vol. 52, no. 4, pp. 65–76, Apr. 2009.
- [29] A. N. Aous, “OpenJPH: An open-source implementation of High-Throughput JPEG 2000,” <https://github.com/aous72/OpenJPH>, 2024.
- [30] Team CharLS, “CharLS: A C/C++ JPEG-LS library implementation,” <https://github.com/team-charls/charls>, 2024.
- [31] NEMA Working Group 4 (WG-04), “DICOM Compression Sample Images,” National Electrical Manufacturers Association, <https://www.dicomstandard.org/Resources/Open-Content/Compression-Samples>, 2024.

- [32] D. A. Clunie, “Breast Tomosynthesis Case 1,” DICOM sample images, <http://www.dclunie.com/pixelmedimagecases/tomo/case1/>, 2009.
- [33] R. Geldreich, “Quick survey of the lossless decompression Pareto frontier,” 2015, <http://richg42.blogspot.com/2015/11/the-lossless-decompression-pareto.html>.
- [34] J. Sneyers, “JPEG XL and the Pareto front,” Cloudinary Blog, 2022, <https://cloudinary.com/blog/jpeg-xl-and-the-pareto-front>.
- [35] J. Tao *et al.*, “FCBench: Cross-domain benchmarking of lossless compression,” *Proc. VLDB Endowment*, vol. 17, no. 7, pp. 1418–1431, 2024.
- [36] A. Anonymous, “Challenges and solutions in selecting optimal lossless data compression algorithms,” arXiv:2509.25219, 2025.
- [37] T. L. Saaty, *The Analytic Hierarchy Process: Planning, Priority Setting, Resource Allocation*. New York, NY, USA: McGraw-Hill, 1980.
- [38] D. A. Clunie, “DICOM support for compression: More than JPEG,” *Medical Imaging Informatics and Teleradiology (MIIT)*, 2009, https://www.dclunie.com/papers/MIIT_2009_DicomCompression_Clunie.pdf.
- [39] K. Sayood, *Introduction to Data Compression*, 5th ed. Burlington, MA, USA: Morgan Kaufmann, 2017.